



Time Series Analysis and Forecasting

Seminar 4: SARIMA Models



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFIN – Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Seminar Outline

Today's Activities:

- 1. Review Quiz** — Checking understanding of SARIMA concepts
- 2. True/False Questions** — Conceptual checks
- 3. Practice Problems** — Calculations with SARIMA
- 4. Worked Examples** — Real-world applications
- 5. Real Data Analysis** — Case study with airline data
- 6. Discussion Topics** — Practical applications
- 7. Summary** — Key takeaways
- 8. AI-Assisted Exercises** — Human vs. AI modeling
- 9. Key Formulas** — Essential SARIMA formulas



Quiz 1: Seasonal Differencing

Question

For monthly data with annual seasonality, what does the operator $(1 - L^{12})$ do?

- A) Takes 12 consecutive differences
- B) Computes $Y_t - Y_{t-12}$
- C) Averages over 12 months
- D) Removes the first 12 observations



Quiz 1: Seasonal Differencing

Question

For monthly data with annual seasonality, what does the operator $(1 - L^{12})$ do?

- A) Takes 12 consecutive differences
- B) Computes $Y_t - Y_{t-12}$
- C) Averages over 12 months
- D) Removes the first 12 observations

Answer: B – Computes $Y_t - Y_{t-12}$

Seasonal difference operator: $(1 - L^{12})Y_t = Y_t - L^{12}Y_t = Y_t - Y_{t-12}$

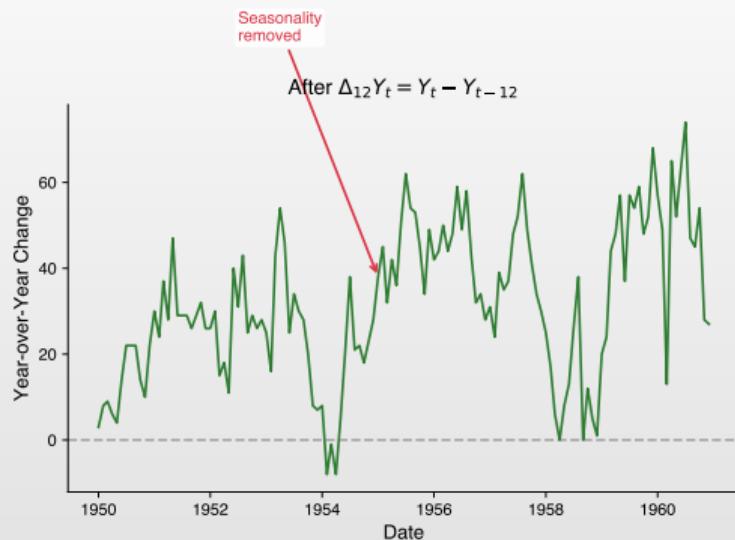
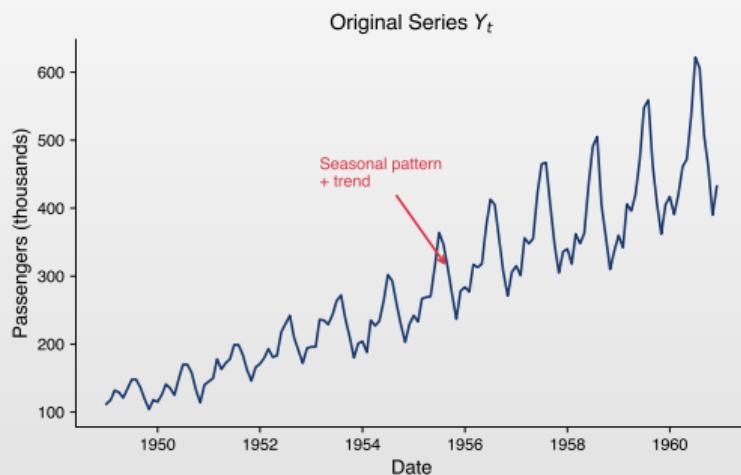
Example (January sales): $Y_{Jan2025} - Y_{Jan2024}$

Effect: Removes stable annual seasonal pattern

Note: $(1 - L^s)$ for any seasonal period s (quarterly: $s = 4$, weekly: $s = 52$)



Visual: Seasonal Difference



Seasonal differencing removes annual patterns by comparing same periods across years.

Q TSA_ch4_def_seasonal_diff



Quiz 2: SARIMA Notation

Question

What does $\text{SARIMA}(1, 1, 1) \times (1, 1, 1)_{12}$ represent?

- A) 12 different ARIMA models
- B) ARIMA with 12 AR and 12 MA terms
- C) ARIMA(1,1,1) with seasonal ARIMA(1,1,1) at period 12
- D) A model requiring 12 years of data



Quiz 2: Answer

Answer: C – ARIMA(1,1,1) with seasonal ARIMA(1,1,1) at period 12

SARIMA($p, d, q \times (P, D, Q)_s$) Notation

$$\phi(L)\Phi(L^s)(1-L)^d(1-L^s)^D Y_t = \theta(L)\Theta(L^s)\varepsilon_t$$

Regular (Non-Seasonal)			Seasonal		
p	= AR order	(Number of AR lags)	P	= Seasonal AR	(SAR lags at s, 2s, ...)
d	= Differencing	(Regular differences)	D	= Seasonal Diff	((1 - L^s) D)
q	= MA order	(Number of MA lags)	Q	= Seasonal MA	(SMA lags at s, 2s, ...)
			s	= Period	(Seasonal period)

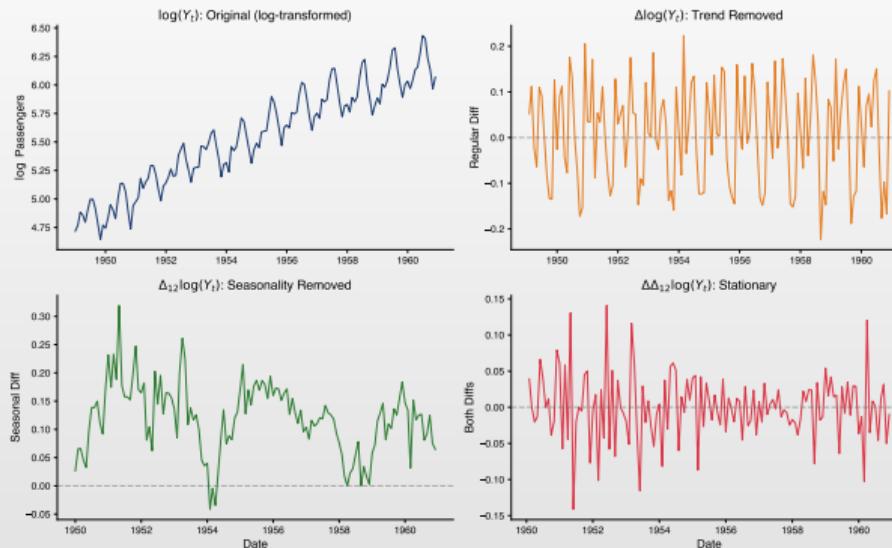
Example: SARIMA(1, 1, 1) \times (0, 1, 1)₁₂

Monthly data with: AR(1), MA(1), one regular diff,
one seasonal diff at lag 12, seasonal MA(1)

$$(1 - \phi_1 L)(1 - \Phi_1 L^{12})(1 - L)(1 - L^{12}) Y_t = (1 + \theta_1 L)(1 + \Theta_1 L^{12}) \varepsilon_t$$



Visual: SARIMA Model Structure



SARIMA combines regular ARIMA components with seasonal components at lag s .

Q TSA_ch4_def_sarima

Quiz 3: The Airline Model

Question

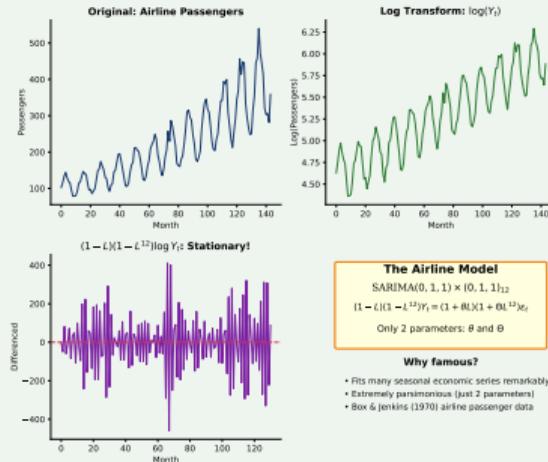
The “airline model” refers to $\text{SARIMA}(0, 1, 1) \times (0, 1, 1)_{12}$. How many parameters does it have (excluding variance)?

- A) 2 parameters
- B) 4 parameters
- C) 6 parameters
- D) 12 parameters



Quiz 3: Answer

Answer: A – 2 parameters (θ_1 and Θ_1)



$$\text{Airline model: } (1 - L)(1 - L^{12})Y_t = (1 + \theta_1L)(1 + \Theta_1L^{12})\varepsilon_t$$

Remarkably fits many seasonal economic series (Box & Jenkins, 1970)



Quiz 4: ACF of Seasonal Data

Question

For monthly data with strong seasonality, where would you expect to see significant ACF spikes?

- A) Only at lag 1
- B) Only at lag 12
- C) At lags 12, 24, 36, ...
- D) Randomly distributed



Quiz 4: ACF of Seasonal Data

Question

For monthly data with strong seasonality, where would you expect to see significant ACF spikes?

- A) Only at lag 1
- B) Only at lag 12
- C) At lags 12, 24, 36, ...
- D) Randomly distributed

Answer: C – At lags 12, 24, 36, ...

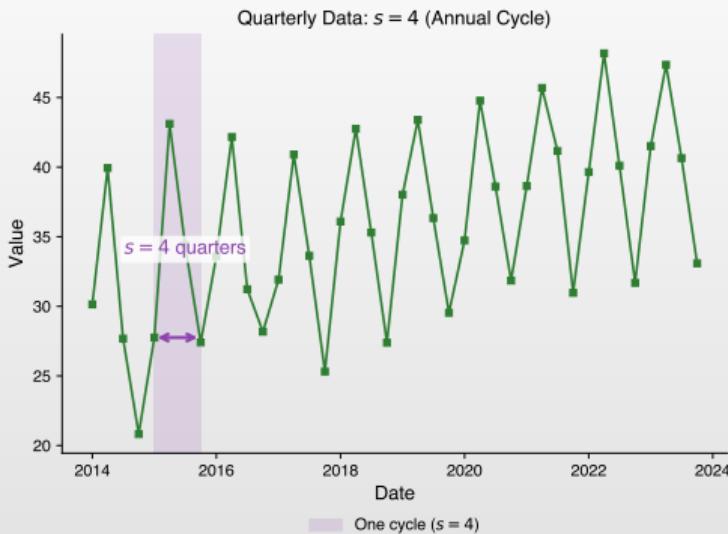
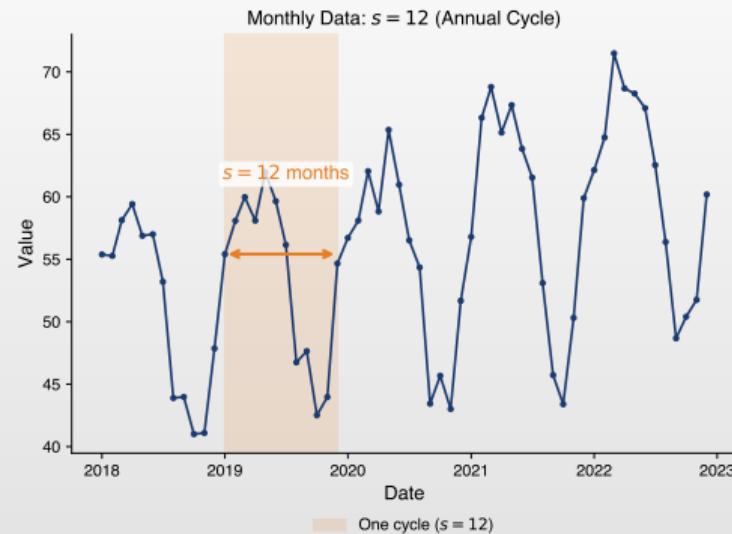
Intuition: January 2024 is similar to January 2023, 2022, etc.

ACF pattern: Spikes at lags $s, 2s, 3s, \dots$ ($\rho_{12}, \rho_{24}, \rho_{36} \neq 0$)

Diagnostic: Slow decay at seasonal lags $\Rightarrow D = 1$; Cutoff after lag $s \Rightarrow Q = 1$



Visual: Seasonality Patterns



Seasonal patterns repeat at regular intervals (monthly, quarterly, etc.) and may be additive or multiplicative.

[TSA_ch4_def_seasonality](#)



Quiz 5: Multiplicative Structure

Question

In SARIMA, what does “multiplicative structure” mean?

- A) The seasonal amplitude grows proportionally
- B) Regular and seasonal polynomials are multiplied
- C) We multiply the data by seasonal factors
- D) The model is estimated using multiplication



Quiz 5: Multiplicative Structure

Question

In SARIMA, what does “multiplicative structure” mean?

- A) The seasonal amplitude grows proportionally
- B) Regular and seasonal polynomials are multiplied
- C) We multiply the data by seasonal factors
- D) The model is estimated using multiplication

Answer: B – Regular and seasonal polynomials are multiplied

Multiplicative SARIMA: $\phi(L)\Phi(L^s)(1-L)^d(1-L^s)^D Y_t = \theta(L)\Theta(L^s)\varepsilon_t$

Example: $(1 - \phi_1 L)(1 - \Phi_1 L^{12}) = 1 - \phi_1 L - \Phi_1 L^{12} + \phi_1 \Phi_1 L^{13}$

Cross-term $\phi_1 \Phi_1 L^{13}$: Captures interaction between short and long dynamics



Quiz 6: Seasonal vs Regular Differencing

Question

When would you apply both regular ($d = 1$) and seasonal ($D = 1$) differencing?

- A) When data has only a trend
- B) When data has only seasonality
- C) When data has both trend and seasonal non-stationarity
- D) Never – they cancel each other



Quiz 6: Seasonal vs Regular Differencing

Question

When would you apply both regular ($d = 1$) and seasonal ($D = 1$) differencing?

- A) When data has only a trend
- B) When data has only seasonality
- C) When data has both trend and seasonal non-stationarity
- D) Never – they cancel each other

Answer: C – Both trend and seasonal non-stationarity

Combined: $W_t = (1 - L)(1 - L^{12})Y_t = Y_t - Y_{t-1} - Y_{t-12} + Y_{t-13}$

When needed: ACF slow decay at lags 1,2,3... $\Rightarrow d = 1$; at lags 12,24,36... $\Rightarrow D = 1$

Examples: Airline passengers, retail sales, energy demand



Quiz 7: Detecting Seasonality from ACF

Question

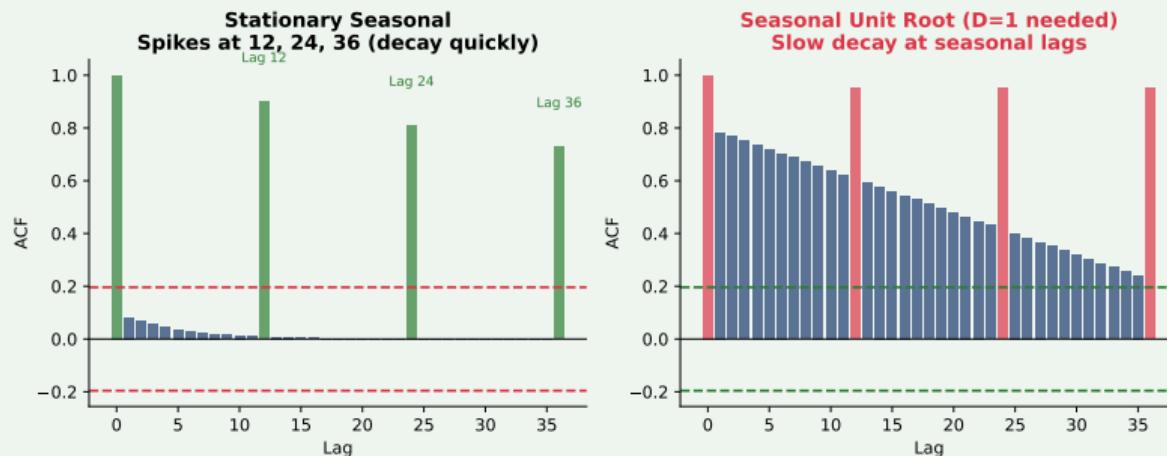
The ACF of a monthly time series shows slow decay at lags 12, 24, and 36. What does this suggest?

- A) The series is stationary
- B) The series needs regular differencing only
- C) The series has a seasonal unit root requiring $D = 1$
- D) The series is white noise



Quiz 7: Answer

Answer: C – Seasonal unit root requiring $D = 1$



Left: Stationary seasonal (fast decay at seasonal lags)
Right: Seasonal unit root (slow decay \Rightarrow need $D = 1$)



Quiz 8: Multiplicative vs Additive Seasonality

Question

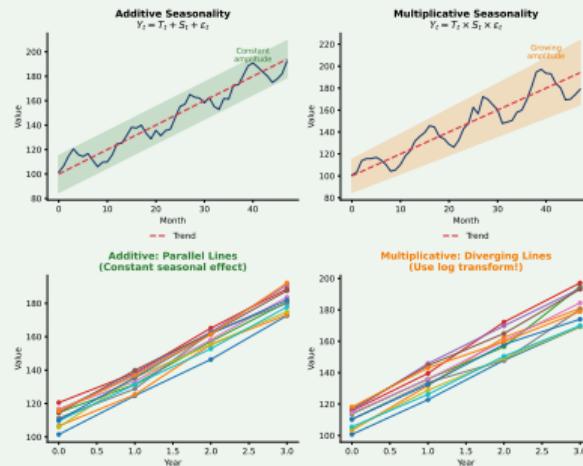
If the seasonal amplitude of a time series grows proportionally with the level, this indicates:

- A) Additive seasonality – use $(1 - L^s)$
- B) Multiplicative seasonality – use log transformation
- C) No seasonality present
- D) Need for regular differencing only



Quiz 8: Answer

Answer: B – Multiplicative seasonality, use log transformation



Multiplicative: Seasonal amplitude grows with level (diverging lines)

Solution: Apply log transformation before fitting SARIMA



Quiz 9: Seasonal Subseries Plot

Question

In a seasonal subseries plot, what indicates multiplicative seasonality?

- A) Lines for each month are parallel
- B) Lines for each month diverge (spread increases over time)
- C) All months have the same mean
- D) Lines are horizontal



Quiz 9: Seasonal Subseries Plot

Question

In a seasonal subseries plot, what indicates multiplicative seasonality?

- A) Lines for each month are parallel
- B) Lines for each month diverge (spread increases over time)
- C) All months have the same mean
- D) Lines are horizontal

Answer: B – Lines diverge (spread increases over time)

Subseries plot: Groups data by month, plots each month's values across years

Parallel \Rightarrow Additive; **Diverging** \Rightarrow Multiplicative; **Horizontal** \Rightarrow No trend

Action: If multiplicative, apply log before fitting SARIMA



Quiz 10: Invertibility in SARIMA

Question

For SARIMA(0, 1, 1) \times (0, 1, 1)₁₂ to be invertible, which condition must hold?

- A) $|\theta_1| < 1$ only
- B) $|\Theta_1| < 1$ only
- C) Both $|\theta_1| < 1$ and $|\Theta_1| < 1$
- D) No invertibility condition exists for MA models



Quiz 10: Invertibility in SARIMA

Question

For SARIMA(0, 1, 1) \times (0, 1, 1)₁₂ to be invertible, which condition must hold?

- A) $|\theta_1| < 1$ only
- B) $|\Theta_1| < 1$ only
- C) Both $|\theta_1| < 1$ and $|\Theta_1| < 1$
- D) No invertibility condition exists for MA models

Answer: C – Both $|\theta_1| < 1$ and $|\Theta_1| < 1$

Invertibility: All MA roots outside unit circle

Multiplicative MA: $(1 + \theta_1 L)(1 + \Theta_1 L^{12})$

Roots: Regular $|z| = |-1/\theta_1| > 1 \Leftrightarrow |\theta_1| < 1$; Seasonal $|\Theta_1| < 1$

Both conditions required for overall invertibility!



Quiz 11: HEGY Test

Question

The HEGY test is used to:

- A) Estimate SARIMA parameters
- B) Test for unit roots at different frequencies (trend and seasonal)
- C) Check residual normality
- D) Compare SARIMA models using information criteria



Quiz 11: HEGY Test

Question

The HEGY test is used to:

- A) Estimate SARIMA parameters
- B) Test for unit roots at different frequencies (trend and seasonal)
- C) Check residual normality
- D) Compare SARIMA models using information criteria

Answer: B – Test for unit roots at different frequencies

HEGY test (Hylleberg-Engle-Granger-Yoo, 1990):

Tests at: Zero freq ($\omega = 0$) $\Rightarrow d = 1$; Nyquist ($\omega = \pi$); Seasonal $\Rightarrow D = 1$

Decision: Reject all \Rightarrow seasonal dummies; Don't reject seasonal \Rightarrow seasonal differencing



Quiz 12: Seasonal MA Identification

Question

After applying $(1 - L)(1 - L^{12})$, the ACF shows a single significant spike at lag 12 only (no spike at lag 1). The PACF decays at seasonal lags. This suggests:

- A) SARIMA(0, 1, 0) \times (0, 1, 1)₁₂
- B) SARIMA(0, 1, 1) \times (0, 1, 0)₁₂
- C) SARIMA(1, 1, 0) \times (1, 1, 0)₁₂
- D) SARIMA(0, 1, 1) \times (0, 1, 1)₁₂



Quiz 12: Seasonal MA Identification

Question

After applying $(1 - L)(1 - L^{12})$, the ACF shows a single significant spike at lag 12 only (no spike at lag 1). The PACF decays at seasonal lags. This suggests:

- A) SARIMA(0, 1, 0) \times (0, 1, 1)₁₂
- B) SARIMA(0, 1, 1) \times (0, 1, 0)₁₂
- C) SARIMA(1, 1, 0) \times (1, 1, 0)₁₂
- D) SARIMA(0, 1, 1) \times (0, 1, 1)₁₂

Answer: A – SARIMA(0, 1, 0) \times (0, 1, 1)₁₂

Pattern: Regular lags – no spikes in ACF/PACF; Seasonal lags – ACF cuts off at s , PACF decays

Interpretation: No regular MA ($q = 0$); Seasonal MA(1) indicated ($Q = 1$)

Model: $(1 - L)(1 - L^{12})Y_t = (1 + \Theta_1 L^{12})\varepsilon_t$



Quiz 13: Over-differencing

Question

After differencing, the ACF shows a large negative spike at lag 1 or lag s . This typically indicates:

- A) The model needs more AR terms
- B) The series has been over-differenced
- C) The series is perfectly stationary
- D) Heteroskedasticity is present



Quiz 13: Over-differencing

Question

After differencing, the ACF shows a large negative spike at lag 1 or lag s . This typically indicates:

- A) The model needs more AR terms
- B) The series has been over-differenced
- C) The series is perfectly stationary
- D) Heteroskedasticity is present

Answer: B – The series has been over-differenced

Signature: ACF at lag 1 $\approx -0.5 \Rightarrow$ over-diff at d ; ACF at lag $s \approx -0.5 \Rightarrow$ over-diff at D

Why? $\Delta^2 Y_t = \varepsilon_t - \varepsilon_{t-1}$ is MA(1) with $\theta = -1$, giving $\rho_1 = -0.5$

Fix: Reduce d or D by one and re-examine ACF/PACF



Quiz 14: Forecasting Horizon

Question

For a SARIMA model with $D = 1$, what happens to forecast confidence intervals as the horizon $h \rightarrow \infty$?

- A) They converge to a fixed width
- B) They grow without bound
- C) They shrink to zero
- D) They oscillate seasonally



Quiz 14: Forecasting Horizon

Question

For a SARIMA model with $D = 1$, what happens to forecast confidence intervals as the horizon $h \rightarrow \infty$?

- A) They converge to a fixed width
- B) They grow without bound
- C) They shrink to zero
- D) They oscillate seasonally

Answer: B – They grow without bound

Unit root property: Any unit root causes unbounded forecast variance

For SARIMA with $D = 1$: $\text{Var}(\hat{Y}_{T+h} - Y_{T+h}) \rightarrow \infty$ as $h \rightarrow \infty$

Intuition: Seasonal shocks accumulate; long-range forecasts have wide CIs



Quiz 15: Seasonal Period Selection

Question

You have daily data showing clear weekly patterns. What seasonal period s should you use in a SARIMA model?

- A) $s = 12$ (monthly)
- B) $s = 7$ (weekly)
- C) $s = 365$ (yearly)
- D) $s = 24$ (hourly)



Quiz 15: Seasonal Period Selection

Question

You have daily data showing clear weekly patterns. What seasonal period s should you use in a SARIMA model?

- A) $s = 12$ (monthly)
- B) $s = 7$ (weekly)
- C) $s = 365$ (yearly)
- D) $s = 24$ (hourly)

Answer: B – $s = 7$ (weekly)

Data	Pattern	Period s
Daily	Weekly	7
Monthly	Annual	12
Quarterly	Annual	4

Rule: $s = \text{observations per cycle of dominant pattern}$



Quiz 16: Seasonal AR Component

Question

In the seasonal component $\Phi(L^s) = 1 - \Phi_1 L^s$, what does the coefficient $\Phi_1 = 0.8$ tell us?

- A) 80% of this period's value comes from the previous period
- B) There is 80% correlation between consecutive observations
- C) 80% of this period's value is explained by the same period last year
- D) The seasonal pattern explains 80% of variance



Quiz 16: Seasonal AR Component

Question

In the seasonal component $\Phi(L^s) = 1 - \Phi_1 L^s$, what does the coefficient $\Phi_1 = 0.8$ tell us?

- A) 80% of this period's value comes from the previous period
- B) There is 80% correlation between consecutive observations
- C) 80% of this period's value is explained by the same period last year
- D) The seasonal pattern explains 80% of variance

Answer: C – 80% explained by same period last year

SAR(1): $Y_t = \Phi_1 Y_{t-12} + \varepsilon_t$

With $\Phi_1 = 0.8$: $Y_{Jan2024} = 0.8 \cdot Y_{Jan2023} + \varepsilon_t$

Interpretation: Strong seasonal persistence – 80% explained by same month last year

Stationarity: Requires $|\Phi_1| < 1$ (satisfied here)



Quiz 17: Seasonal Stationarity

Question

A seasonal process with $\Phi_1 = 1$ in $\text{SARIMA}(0, 0, 0) \times (1, 0, 0)_{12}$ is:

- A) Stationary
- B) Has a seasonal unit root (seasonally integrated)
- C) Explosive
- D) Undefined



Quiz 17: Seasonal Stationarity

Question

A seasonal process with $\Phi_1 = 1$ in $\text{SARIMA}(0, 0, 0) \times (1, 0, 0)_{12}$ is:

- A) Stationary
- B) Has a seasonal unit root (seasonally integrated)
- C) Explosive
- D) Undefined

Answer: B – Has a seasonal unit root

Model: $Y_t = Y_{t-12} + \varepsilon_t$ (seasonal random walk)

Properties: Variance grows with time; each month follows its own RW; need $D = 1$

Analogy: Like regular random walk but at seasonal frequency



Quiz 18: Model Comparison

Question

Model A: SARIMA(1, 1, 1) \times (1, 1, 1)₁₂ has AIC = 520. Model B: SARIMA(0, 1, 1) \times (0, 1, 1)₁₂ has AIC = 525. Which statement is most accurate?

- A) Model A is always better since it has lower AIC
- B) Model B should be preferred due to parsimony despite higher AIC
- C) The AIC difference of 5 suggests Model A is substantially better
- D) We cannot compare models with different orders



Quiz 18: Model Comparison

Question

Model A: SARIMA(1, 1, 1) \times (1, 1, 1)₁₂ has AIC = 520. Model B: SARIMA(0, 1, 1) \times (0, 1, 1)₁₂ has AIC = 525. Which statement is most accurate?

- A) Model A is always better since it has lower AIC
- B) Model B should be preferred due to parsimony despite higher AIC
- C) The AIC difference of 5 suggests Model A is substantially better
- D) We cannot compare models with different orders

Answer: C – AIC difference of 5 suggests Model A is substantially better

Rule of thumb: $\Delta\text{AIC} < 2$: equivalent; 2–10: some evidence; > 10 : strong evidence

Here: $\Delta\text{AIC} = 5$ suggests Model A meaningfully better

Always: Also check residual diagnostics and forecast performance!



Quiz 19: Seasonal Patterns in Residuals

Question

After fitting a SARIMA model, you notice significant ACF spikes at lags 12 and 24 in the residuals. What does this indicate?

- A) The model is correctly specified
- B) The seasonal component is inadequate
- C) The data is not seasonal
- D) Overfitting has occurred



Quiz 19: Seasonal Patterns in Residuals

Question

After fitting a SARIMA model, you notice significant ACF spikes at lags 12 and 24 in the residuals. What does this indicate?

- A) The model is correctly specified
- B) The seasonal component is inadequate
- C) The data is not seasonal
- D) Overfitting has occurred

Answer: B – The seasonal component is inadequate

Diagnostics: Good residuals should be white noise (no significant ACF)

Seasonal ACF in residuals: Model hasn't captured seasonal structure; try increasing P or Q ; verify D is correct

Action: Try SARIMA with higher seasonal order, check Ljung-Box at seasonal lags



Quiz 20: Practical Forecasting

Question

You're forecasting monthly retail sales with $\text{SARIMA}(0, 1, 1) \times (0, 1, 1)_{12}$. For the 13-month-ahead forecast, which historical observations are most influential?

- A) Only the most recent observation
- B) The observation from the same month last year
- C) All observations equally
- D) Only observations from the same month in all previous years



Quiz 20: Practical Forecasting

Question

You're forecasting monthly retail sales with SARIMA(0, 1, 1) \times (0, 1, 1)₁₂. For the 13-month-ahead forecast, which historical observations are most influential?

- A) Only the most recent observation
- B) The observation from the same month last year
- C) All observations equally
- D) Only observations from the same month in all previous years

Answer: B – The observation from the same month last year

For 13-month ahead: Most influential is Y_{T-11} (same month last year), also Y_T and Y_{T-12}

Intuition: "Next January looks like last January, adjusted for recent trend"



True/False Questions (1-6)

Question

Determine whether each statement is True or False:

1. The seasonal period s for quarterly data with annual patterns is $s = 4$.
2. SARIMA models can only handle one seasonal frequency.
3. If AIC selects SARIMA(1,1,1) \times (1,1,1)₁₂ and BIC selects the airline model, BIC is always wrong.
4. The Kruskal-Wallis test can detect seasonality without assuming normality.
5. After fitting a SARIMA model, residuals should show no significant ACF at seasonal lags.
6. Log transformation converts multiplicative seasonality to additive.

Answer on next slide...



True/False Solutions (1-6)

Answers

1. **TRUE:** Quarterly data with annual cycle has $s = 4$ quarters per year.
2. **TRUE:** Standard SARIMA handles one s ; multiple seasonalities need TBATS or Fourier terms.
3. **FALSE:** BIC penalizes complexity more; simpler model may be better for interpretation/forecasting.
4. **TRUE:** Kruskal-Wallis is nonparametric, comparing distributions across seasons.
5. **TRUE:** Residual ACF should be within confidence bands at ALL lags including seasonal.
6. **TRUE:** $\log(T \times S \times \varepsilon) = \log T + \log S + \log \varepsilon$ (additive form).



Problem 1: Expanding the Seasonal Difference

Exercise

Expand $(1 - L)(1 - L^{12})Y_t$ fully. What observations are involved?



Problem 1: Expanding the Seasonal Difference

Exercise

Expand $(1 - L)(1 - L^{12})Y_t$ fully. What observations are involved?

Solution

$$(1 - L)(1 - L^{12}) = 1 - L - L^{12} + L^{13}$$

$$\text{Therefore: } (1 - L)(1 - L^{12})Y_t = Y_t - Y_{t-1} - Y_{t-12} + Y_{t-13}$$

Interpretation: This is the difference of differences:

- First seasonal difference: $Y_t - Y_{t-12}$ (this year vs last year)
- Then regular difference: $(Y_t - Y_{t-12}) - (Y_{t-1} - Y_{t-13})$



Problem 2: Airline Model Expansion

Exercise

Write out the full equation for the airline model SARIMA(0, 1, 1) \times (0, 1, 1)₁₂:

$$(1 - L)(1 - L^{12})Y_t = (1 + \theta_1 L)(1 + \Theta_1 L^{12})\varepsilon_t$$



Problem 2: Airline Model Expansion

Exercise

Write out the full equation for the airline model SARIMA(0, 1, 1) \times (0, 1, 1)₁₂:

$$(1 - L)(1 - L^{12})Y_t = (1 + \theta_1 L)(1 + \Theta_1 L^{12})\varepsilon_t$$

Solution

Expand the MA side: $(1 + \theta_1 L)(1 + \Theta_1 L^{12}) = 1 + \theta_1 L + \Theta_1 L^{12} + \theta_1 \Theta_1 L^{13}$

Full model: $Y_t - Y_{t-1} - Y_{t-12} + Y_{t-13} = \varepsilon_t + \theta_1 \varepsilon_{t-1} + \Theta_1 \varepsilon_{t-12} + \theta_1 \Theta_1 \varepsilon_{t-13}$

Note: The cross-term $\theta_1 \Theta_1 L^{13}$ is the multiplicative interaction between regular and seasonal MA components.



Problem 3: Parameter Count

Exercise

How many parameters (excluding σ^2) are in SARIMA(2, 1, 1) \times (1, 0, 1)₄?



Problem 3: Parameter Count

Exercise

How many parameters (excluding σ^2) are in SARIMA(2, 1, 1) \times (1, 0, 1)₄?

Solution

- Regular AR($p = 2$): $\phi_1, \phi_2 \Rightarrow 2$ parameters
- Regular MA($q = 1$): $\theta_1 \Rightarrow 1$ parameter
- Seasonal AR($P = 1$): $\Phi_1 \Rightarrow 1$ parameter
- Seasonal MA($Q = 1$): $\Theta_1 \Rightarrow 1$ parameter

Total: 5 parameters

Note: The differencing orders ($d = 1, D = 0$) don't add parameters – they're transformations applied to the data.



Problem 4: SARIMA Forecasting

Exercise

Given the airline model with $\theta_1 = -0.4$ and $\Theta_1 = -0.6$, and:

- $Y_T = 500, Y_{T-1} = 495, Y_{T-11} = 480, Y_{T-12} = 470$
- $\varepsilon_T = 5, \varepsilon_{T-11} = -3, \varepsilon_{T-12} = 2$

Forecast Y_{T+1} .



Problem 4: SARIMA Forecasting

Exercise

Given the airline model with $\theta_1 = -0.4$ and $\Theta_1 = -0.6$, and:

- $Y_T = 500, Y_{T-1} = 495, Y_{T-11} = 480, Y_{T-12} = 470$
- $\varepsilon_T = 5, \varepsilon_{T-11} = -3, \varepsilon_{T-12} = 2$

Forecast Y_{T+1} .

Solution

From the model: $Y_{T+1} = Y_T + Y_{T-11} - Y_{T-12} + \varepsilon_{T+1} + \theta_1\varepsilon_T + \Theta_1\varepsilon_{T-11} + \theta_1\Theta_1\varepsilon_{T-12}$

Setting $\mathbb{E}[\varepsilon_{T+1}] = 0$:

$$\begin{aligned}\hat{Y}_{T+1} &= 500 + 480 - 470 + 0 + (-0.4)(5) + (-0.6)(-3) + (-0.4)(-0.6)(2) \\ &= 510 - 2 + 1.8 + 0.48 = 510.28\end{aligned}$$



Problem 5: Identifying Seasonal Period

Exercise

Match each data type with its typical seasonal period s :

1. Quarterly GDP data
2. Monthly retail sales
3. Weekly restaurant reservations
4. Daily electricity demand



Problem 5: Identifying Seasonal Period

Exercise

Match each data type with its typical seasonal period s :

1. Quarterly GDP data
2. Monthly retail sales
3. Weekly restaurant reservations
4. Daily electricity demand

Solution

1. Quarterly GDP: $s = 4$ (annual cycle over 4 quarters)
2. Monthly retail sales: $s = 12$ (annual cycle over 12 months)
3. Weekly restaurant reservations: $s = 7$ (weekly cycle) or $s = 52$ (annual)
4. Daily electricity demand: $s = 7$ (weekly pattern) or $s = 365$ (annual)

Note: Some series have multiple seasonal patterns (e.g., daily data may have weekly AND annual cycles).



Example: Monthly Retail Sales Analysis

Scenario

You have 5 years of monthly retail sales data showing clear December peaks and January troughs. Build an appropriate SARIMA model.

Step-by-step Approach

1. **Visual inspection:** Plot shows upward trend + strong December spikes
2. **Seasonal period:** Monthly data with annual pattern $\Rightarrow s = 12$
3. **Transformation:** Consider $\log(Y_t)$ if seasonal amplitude grows with level
4. **Differencing:** Try $(1 - L)(1 - L^{12})Y_t$ – check ACF/PACF
5. **Model selection:** Start with airline model, compare via AIC



Example: ACF/PACF Interpretation for Seasonal Data

Observed Patterns (after differencing)

- ACF: Significant at lags 1, 12; cuts off after lag 1 and lag 12
- PACF: Significant at lags 1, 12, 13; decays at multiples of 12

Interpretation

Regular component: ACF cuts off at 1 \Rightarrow MA(1)

Seasonal component: ACF significant only at lag 12 \Rightarrow seasonal MA(1)

Suggested model: SARIMA(0, d, 1) \times (0, D, 1)₁₂ – the airline model!

Alternative check: If PACF showed cutoff at seasonal lags instead of ACF, consider seasonal AR terms.



Example: Python Implementation

Fitting SARIMA in Python

```
from statsmodels.tsa.statespace.sarimax import SARIMAX
import pmdarima as pm

# Manual fit
model = SARIMAX(y, order=(0,1,1), seasonal_order=(0,1,1,12))
results = model.fit()
print(results.summary())

# Automatic selection
auto_model = pm.auto_arima(y, seasonal=True, m=12,
                            start_p=0, max_p=2,
                            start_q=0, max_q=2,
                            d=1, D=1,
                            trace=True)
```



Example: Interpreting SARIMA Output

Sample statsmodels Output

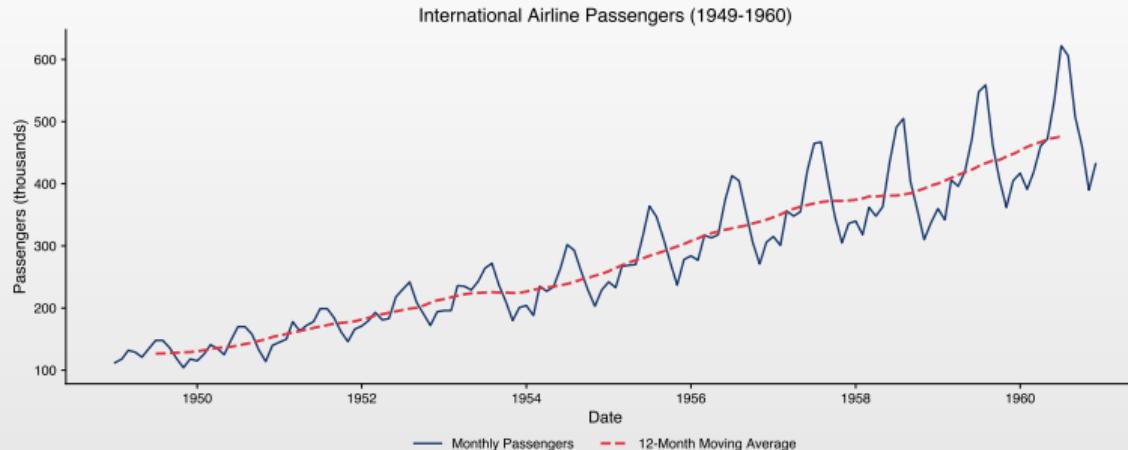
```
SARIMAX Results
=====
Model:           SARIMAX(0,1,1)x(0,1,1,12)   AIC:    1348.52
                           BIC:    1358.21
=====
                  coef      std err      z     P>|z|
-----
ma.L1          -0.4018      0.072    -5.58    0.000
ma.S.L12        -0.5521      0.081    -6.82    0.000
sigma2         1254.3201    142.856     8.78    0.000
```

Interpretation

- $\hat{\theta}_1 = -0.40$: Negative MA means positive shocks reduce next period's value
- $\hat{\Theta}_1 = -0.55$: Same-season correlation is captured
- Both coefficients significant ($p < 0.001$); $|\theta|, |\Theta| < 1$ – invertible



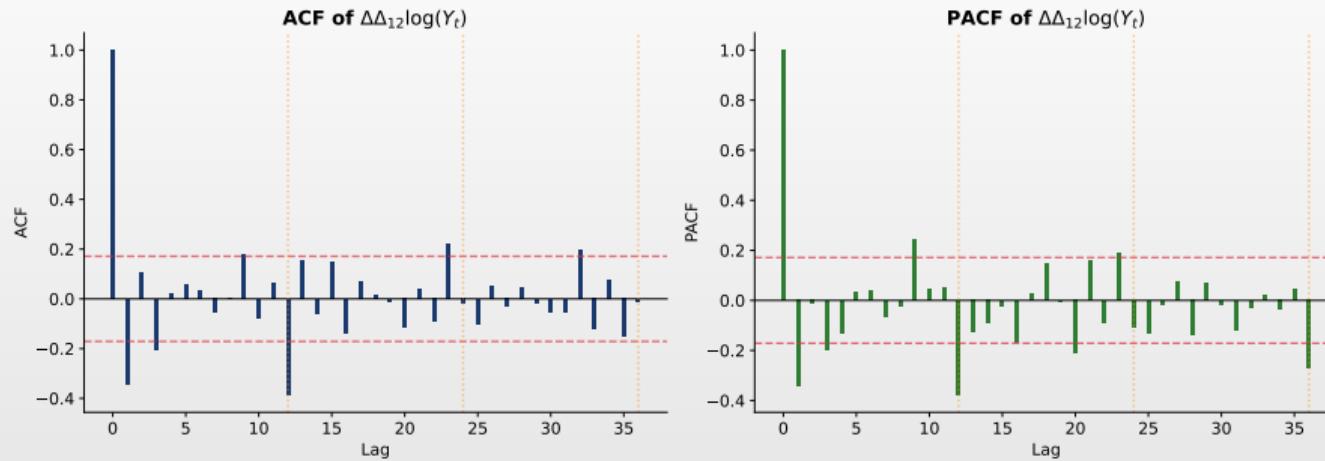
Case Study: Airline Passengers (1949–1960)



- Classic Box-Jenkins dataset: 144 monthly observations
- Clear **upward trend** and **seasonal pattern** (summer peaks)
- Seasonal amplitude **grows with level** ⇒ multiplicative seasonality
- Suggests: log transformation + SARIMA modeling



ACF/PACF Analysis After Differencing



- ◻ After $(1 - L)(1 - L^{12}) \log(Y_t)$: series appears stationary
- ◻ Significant spike at lag 1 in ACF \Rightarrow MA(1) component
- ◻ Significant spike at lag 12 in ACF \Rightarrow Seasonal MA(1) component
- ◻ Pattern suggests: **SARIMA(0, 1, 1)(0, 1, 1)₁₂** (airline model)



SARIMA Estimation Results: Airline Data

Model: SARIMA(0, 1, 1)(0, 1, 1)₁₂ on log(Passengers)

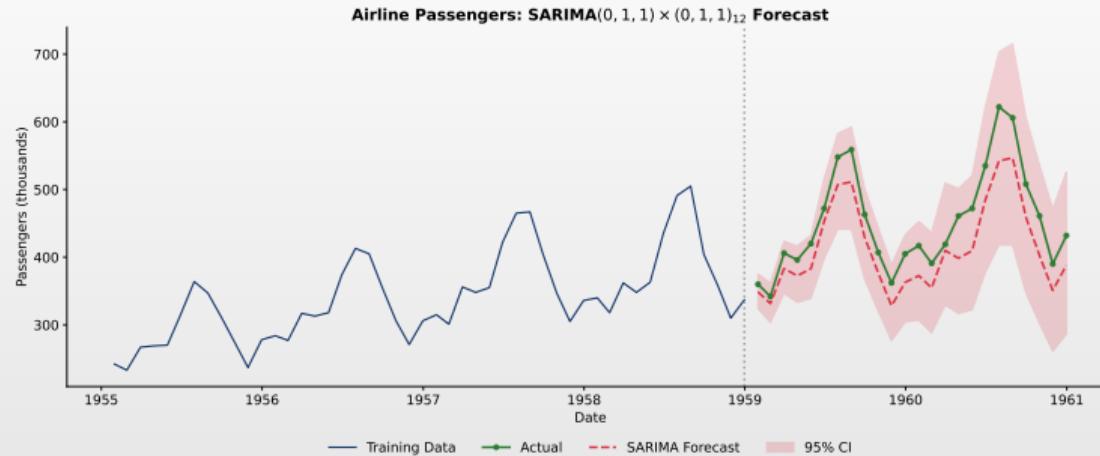
Parameter	Estimate	Std. Error	z-stat	p-value
θ_1 (MA.L1)	-0.4018	0.0896	-4.48	< 0.001
Θ_1 (MA.S.L12)	-0.5569	0.0731	-7.62	< 0.001
σ^2	0.00135	-	-	-

Model Fit Statistics

- Log-Likelihood: 244.70
- AIC: -483.40, BIC: -474.53
- Both MA coefficients significant and within invertibility bounds



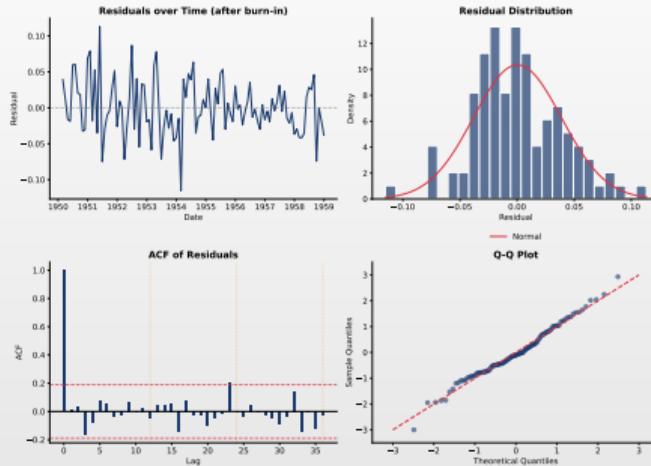
Forecast: 24 Months Ahead



- Forecasts capture both trend and seasonal pattern
- 95% confidence intervals widen over forecast horizon
- Seasonal peaks (July-August) and troughs (February) clearly visible
- Model successfully extrapolates the multiplicative seasonal pattern



Model Diagnostics



- Residuals appear random with no systematic patterns
- Distribution approximately normal (Q-Q plot close to diagonal)
- ACF of residuals within confidence bounds – no significant autocorrelation
- Ljung-Box test: $p > 0.05$ at all tested lags \Rightarrow adequate model



Discussion: Deterministic vs Stochastic Seasonality

Key Question

When should you use seasonal dummies vs SARIMA for seasonal data?

Considerations

Seasonal dummies (deterministic):

- Fixed, repeating pattern each year
- Same December effect every year
- Appropriate when seasonality is stable

SARIMA (stochastic):

- Evolving seasonal pattern
- This year's December depends on last year's December
- Better when seasonal amplitude varies



Discussion: Log Transformation

Key Question

When should you take logarithms before fitting SARIMA?

Guidelines

Use log transformation when:

- Seasonal fluctuations grow with the level (multiplicative seasonality)
- Variance increases over time
- Data is strictly positive (prices, sales, counts)

Avoid log when:

- Seasonal pattern is additive (constant amplitude)
- Data contains zeros or negatives
- Already on a rate/ratio scale

Tip: Compare AIC of models with and without log transformation.



Discussion: Multiple Seasonalities

Challenge

Daily sales data may have both weekly (7-day) and annual (365-day) seasonal patterns. How do you handle this?

Approaches

1. **Nested SARIMA:** Model at shorter frequency, include longer as exogenous
2. **TBATS/BATS models:** Explicitly handle multiple seasonalities
3. **Fourier terms:** Add sin/cos terms for each seasonal frequency
4. **Prophet/similar:** Modern tools designed for multiple seasonalities

Note: Standard SARIMA handles only one seasonal period. For complex seasonality, consider specialized methods.



Discussion: Forecasting Seasonal Data

Key Question

What are the unique challenges of forecasting seasonal time series?

Challenges and Solutions

- Horizon matters:** 12-month forecast means predicting a full cycle
- Uncertainty grows:** Seasonal forecasts compound regular uncertainty
- Turning points:** Capturing when seasons peak/trough
- Structural breaks:** COVID-19 disrupted many seasonal patterns

Best practices:

- Use rolling-origin cross-validation
- Compare against seasonal naive benchmark
- Report forecast intervals, especially at seasonal horizons



Take-Home Exercises

1. **Theoretical:** Show that $(1 - L)(1 - L^4)$ can be written as $(1 - L - L^4 + L^5)$ and explain what this transformation does to quarterly data with annual seasonality.
2. **Computation:** For SARIMA(1, 0, 0) \times (1, 0, 0)₄ with $\phi_1 = 0.5$ and $\Phi_1 = 0.8$, write out the full AR polynomial and identify all non-zero coefficients.
3. **Applied:** Download monthly airline passenger data and:
 - ▶ Plot the series and identify trend/seasonality
 - ▶ Apply appropriate transformations
 - ▶ Fit the airline model and interpret coefficients
 - ▶ Generate 24-month forecasts with confidence intervals
4. **Comparison:** Fit both SARIMA(0, 1, 1) \times (0, 1, 1)₁₂ and SARIMA(1, 1, 0) \times (1, 1, 0)₁₂ to the airline data. Compare using AIC, BIC, and residual diagnostics. Which is preferred?



Exercise Solutions Hints

Hints

1. Expand $(1 - L)(1 - L^4) = 1 \cdot 1 - 1 \cdot L^4 - L \cdot 1 + L \cdot L^4 = 1 - L - L^4 + L^5$
2. AR polynomial: $(1 - \phi_1 L)(1 - \Phi_1 L^4) = 1 - 0.5L - 0.8L^4 + 0.4L^5$
3. For airline data:
 - ▶ Use log transformation (multiplicative seasonality)
 - ▶ Both $d = 1$ and $D = 1$ needed
 - ▶ Typical estimates: $\theta_1 \approx -0.4$, $\Theta_1 \approx -0.6$
4. The MA-based airline model typically fits better than pure AR seasonal model for this data (lower AIC).

Key Points from Today's Seminar

What We Covered

1. **Seasonal differencing:** $(1 - L^s)$ removes stochastic seasonality
2. **SARIMA notation:** $(p, d, q) \times (P, D, Q)_s$ separates regular and seasonal
3. **The airline model:** SARIMA(0, 1, 1)(0, 1, 1)₁₂ is surprisingly effective
4. **Multiplicative structure:** Creates interaction terms between components
5. **ACF/PACF diagnostics:** Patterns at both regular and seasonal lags
6. **Log transformation:** Often needed for multiplicative seasonality

Next Seminar

GARCH and Volatility Modeling: volatility clustering, ARCH/GARCH estimation, financial risk applications (VaR).



AI in SARIMA Modeling

Context

AI tools can fit SARIMA models and generate seasonal forecasts automatically. The critical skill is **evaluating whether the seasonal methodology is correct**.

Key questions to ask about any AI-generated SARIMA analysis:

1. Did it check for multiplicative seasonality and apply log if needed?
2. Is the seasonal period s correctly identified from the data frequency?
3. Did it test for both regular and seasonal unit roots before choosing d and D ?
4. Are residuals free of seasonal patterns (check ACF at lags $s, 2s, 3s$)?
5. Did it compare against a seasonal naive benchmark?



AI Exercise 1: Critique an AI SARIMA Analysis

Scenario

You asked an AI: "Fit the best SARIMA model to this monthly retail sales data." It returned:

- Fitted SARIMA(2,1,2)(2,1,2)₁₂ with AIC = 1520
- No check for multiplicative vs additive seasonality
- Applied $D = 1$ "because the data is monthly"
- Ljung-Box p-value = 0.03 on residuals
- 24-month forecast with constant-width confidence intervals

Your critique:

1. Why is the model likely over-parameterized? How many parameters does it have?
2. Why should multiplicative seasonality be checked before fitting?
3. Why is Ljung-Box p = 0.03 **not** acceptable at 5% level?
4. Why must SARIMA confidence intervals widen (not stay constant)?



AI Exercise 2: Prompt Refinement for SARIMA

Task

Iteratively improve prompts for fitting a SARIMA model to monthly airline data.

Round 1 (vague): “*Fit a seasonal model to monthly airline passengers*”

- What did the AI produce? Did it check for multiplicative seasonality?

Round 2 (better): “*Check if seasonality is multiplicative (log transform if so), apply seasonal differencing, fit SARIMA using AIC, check residual ACF at seasonal lags*”

- Did the AI follow the seasonal Box-Jenkins methodology?

Round 3 (expert): “*(1) Plot series, check if amplitude grows with level, (2) if multiplicative, apply log, (3) test ADF on levels & seasonal differences, (4) fit airline model + 2 alternatives, (5) compare AIC/BIC, (6) Ljung-Box on residuals, (7) 24-month rolling forecast with RMSE*”

- Compare results across all three rounds



AI Exercise 3: Model Selection Competition

Task

Download the classic airline passenger dataset (Box-Jenkins, 1970).

Your approach (manual):

- Check multiplicative vs additive seasonality → log if needed
- Determine d and D from ADF/KPSS and seasonal ACF
- ACF/PACF of transformed series → candidate models
- Compare airline model vs SARIMA(1, 1, 0)(1, 1, 0)₁₂ via AIC/BIC
- Rolling 1-step forecast on last 24 observations

AI approach:

- Ask AI to “find the best SARIMA model and forecast airline passengers”

Compare:

- Did the AI apply log transformation? Which model did each select?
- Compare out-of-sample RMSE; did the AI check seasonal residual patterns?
- Submit:** 1-page reflection on AI strengths and weaknesses



Key Formulas Summary

Concept	Formula
Seasonal difference	$(1 - L^s)Y_t = Y_t - Y_{t-s}$
Combined differencing	$(1 - L)(1 - L^s)Y_t = Y_t - Y_{t-1} - Y_{t-s} + Y_{t-s-1}$
SARIMA(p, d, q)(P, D, Q) $_s$	$\phi(L)\Phi(L^s)(1 - L)^d(1 - L^s)^D Y_t = \theta(L)\Theta(L^s)\varepsilon_t$
Airline model	$(1 - L)(1 - L^{12})Y_t = (1 + \theta_1 L)(1 + \Theta_1 L^{12})\varepsilon_t$
Multiplicative AR	$(1 - \phi_1 L)(1 - \Phi_1 L^s) = 1 - \phi_1 L - \Phi_1 L^s + \phi_1 \Phi_1 L^{s+1}$
Multiplicative MA	$(1 + \theta_1 L)(1 + \Theta_1 L^s) = 1 + \theta_1 L + \Theta_1 L^s + \theta_1 \Theta_1 L^{s+1}$
Invertibility	$ \theta_i < 1$ and $ \Theta_j < 1$ for all i, j
Stationarity	$ \Phi_j < 1$ and regular AR roots outside unit circle
Log transform	$\log(T \times S \times \varepsilon) = \log T + \log S + \log \varepsilon$

Notation: s = seasonal period, ϕ/Φ = regular/seasonal AR, θ/Θ = regular/seasonal MA, d/D = regular/seasonal differencing order



What's Next?

Seminar 5: GARCH and Volatility Modeling

- Volatility clustering:** Why financial returns show time-varying variance
- ARCH/GARCH models:** Modeling conditional heteroskedasticity
- Extensions:** EGARCH, GJR-GARCH, leverage effects
- Applications:** Value-at-Risk and risk management

Questions?



Bibliography I

Fundamental textbooks

- Hyndman, R.J., & Athanasopoulos, G. (2021). *Forecasting: Principles and Practice*, 3rd ed., OTexts.
- Shumway, R.H., & Stoffer, D.S. (2017). *Time Series Analysis and Its Applications*, 4th ed., Springer.
- Brockwell, P.J., & Davis, R.A. (2016). *Introduction to Time Series and Forecasting*, 3rd ed., Springer.

Financial time series

- Tsay, R.S. (2010). *Analysis of Financial Time Series*, 3rd ed., Wiley.
- Franke, J., Härdle, W.K., & Hafner, C.M. (2019). *Statistics of Financial Markets*, 4th ed., Springer.



Bibliography II

Modern approaches and Machine Learning

- Nielsen, A. (2019). *Practical Time Series Analysis*, O'Reilly Media.
- Petropoulos, F., et al. (2022). *Forecasting: Theory and Practice*, International Journal of Forecasting.
- Makridakis, S., Spiliotis, E., & Assimakopoulos, V. (2020). The M4 Competition, International Journal of Forecasting.

Online resources and code

- **Quantlet:** <https://quantlet.com> — Code repository for statistics
- **Quantinar:** <https://quantinar.com> — Learning platform for quantitative methods
- **GitHub TSA:** https://github.com/QuantLet/TSA/tree/main/TSA_ch4 — Python code for this seminar

