



Time Series Analysis and Forecasting

Chapter 10: Comprehensive Review

Applied Case Studies with Rigorous Methodology



Outline

- ① Forecasting Methodology
- ② Case Study 1: Bitcoin Volatility (GARCH)
- ③ Case Study 2: Sunspot Cycles (Fourier)
- ④ Case Study 3: Unemployment (Prophet)
- ⑤ Case Study 4: Multivariate Analysis (VAR)
- ⑥ Synthesis and Guidelines

The Scientific Approach to Forecasting

Research Question

How do we **rigorously evaluate** forecast performance while avoiding overfitting?

The Fundamental Problem

- In-sample fit \neq Out-of-sample performance
- Models can “memorize” training data without learning patterns
- **Solution:** Proper train/validation/test methodology

Key Principle

“The test set must remain **untouched** until final evaluation.”

— Standard practice in machine learning and econometrics

Time Series Train/Validation/Test Split



Training Set	Validation Set	Test Set
<ul style="list-style-type: none">• Fit parameters• Largest portion	<ul style="list-style-type: none">• Compare models• Tune hyperparams	<ul style="list-style-type: none">• Held out• Final metrics

Definition 1 (Forecast Error Metrics)

Let y_t be actual, \hat{y}_t forecast:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_t (y_t - \hat{y}_t)^2}, \quad \text{MAE} = \frac{1}{n} \sum_t |y_t - \hat{y}_t|, \quad \text{MAPE} = \frac{100\%}{n} \sum_t \left| \frac{y_t - \hat{y}_t}{y_t} \right| \quad (1)$$

When to Use Each

- **RMSE**: Penalizes large errors
- **MAE**: Robust to outliers
- **MAPE**: Scale-independent (%)

Caution

- MAPE undefined when $y_t = 0$
- Compare on **same** test set
- Report **out-of-sample** metrics

Bitcoin: Problem Statement

Research Question

Can we forecast Bitcoin's **volatility** using GARCH models?

Data Characteristics

- Source: Yahoo Finance (BTC-USD)
- Period: Jan 2019 – Jan 2025
- Frequency: Daily
- Observations: $\approx 2,200$ days

Key Insight

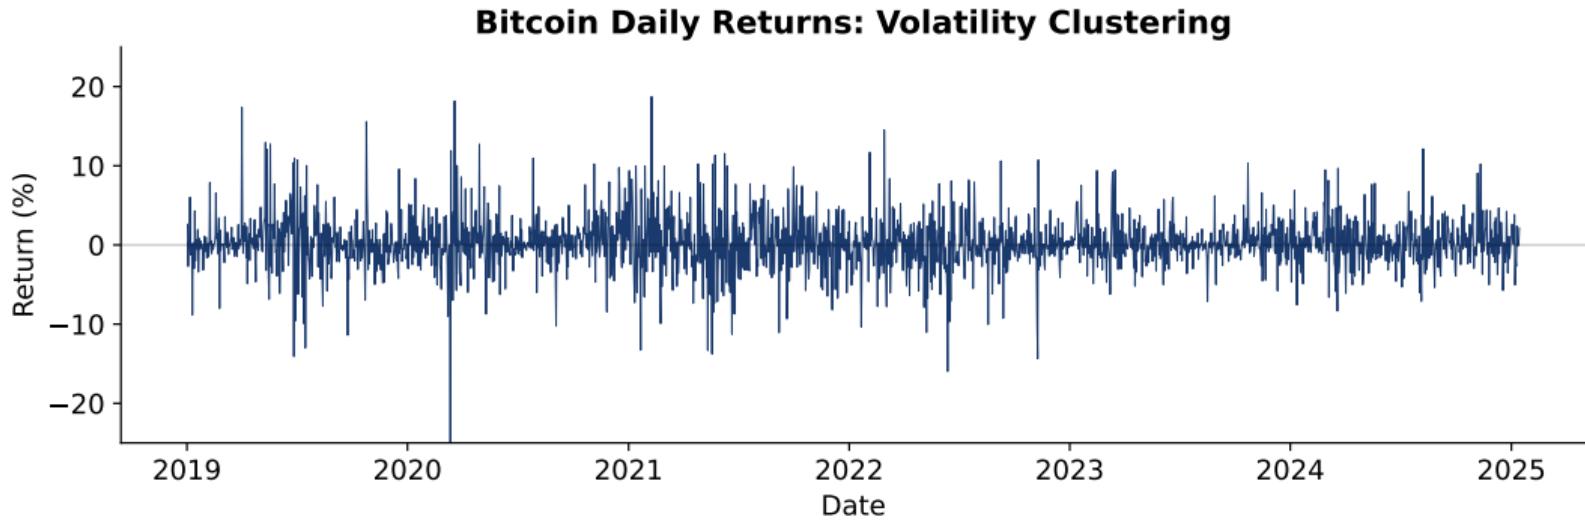
Financial returns are typically:

- **Unpredictable** in mean
- **Predictable** in variance

⇒ Focus on **volatility forecasting**

Stylized Facts

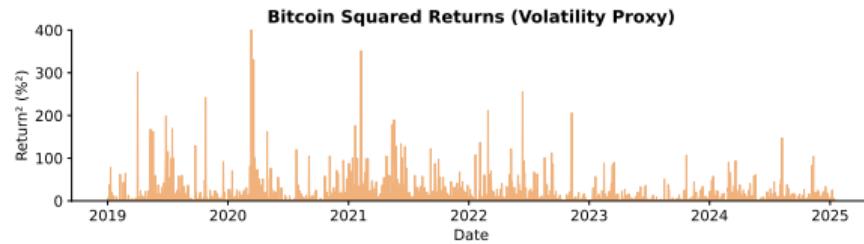
- Returns: near-zero mean
- Fat tails (kurtosis > 3)
- Volatility clustering



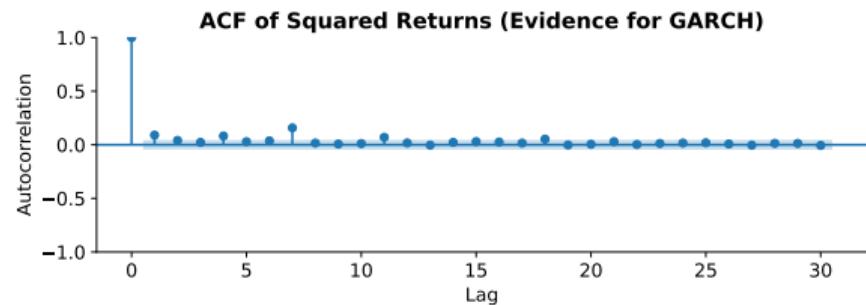
Observation

Large returns tend to follow large returns, small follow small. This is **volatility clustering**—the phenomenon GARCH captures.

Bitcoin: Evidence for GARCH



Squared returns r_t^2 proxy for volatility σ_t^2 . Spikes cluster together.



ACF bars exceed blue bands \Rightarrow significant autocorrelation at multiple lags.

Why GARCH?

If r_t^2 were white noise, ACF would be zero. Significant ACF means **past volatility predicts future volatility**—GARCH captures this!

GARCH Model Specification

Definition 2 (GARCH(p,q) Model)

Let r_t denote returns. The GARCH(p,q) model is:

$$r_t = \mu + \varepsilon_t, \quad \varepsilon_t = \sigma_t z_t, \quad z_t \sim N(0, 1) \quad (2)$$

$$\sigma_t^2 = \omega + \sum_{i=1}^q \alpha_i \varepsilon_{t-i}^2 + \sum_{j=1}^p \beta_j \sigma_{t-j}^2 \quad (3)$$

where $\omega > 0$, $\alpha_i \geq 0$, $\beta_j \geq 0$, and $\sum_{i=1}^q \alpha_i + \sum_{j=1}^p \beta_j < 1$.

Model Variants

- **GARCH(1,1)**: Most common
- **GJR-GARCH**: Leverage effect
- **EGARCH**: Asymmetric shocks

Interpretation

- α : Impact of past shocks
- β : Persistence of volatility
- $\alpha + \beta \approx 1$: High persistence

Bitcoin: Data Split and Stationarity

Data Split

Set	Period	N
Training	2019-01 to 2022-09	1,365
Validation	2022-09 to 2023-10	400
Test	2023-10 to 2025-01	435
Total		2,200

Stationarity Tests

Series	ADF	Result
Prices	$p = 0.50$	Non-stationary
Returns	$p < 0.01$	Stationary

⇒ Model **returns**, not prices

Why Stationarity Matters

GARCH requires weakly stationary input. Prices follow random walk; returns are stationary.

Bitcoin: Model Selection on Validation Set

Methodology

Fit each model on **training data**, evaluate on **validation set**.

Model	AIC	BIC	Val MAE	Selection
GARCH(1,1)	6,994.8	7,020.6	2.638	Best
GARCH(2,1)	6,993.7	7,024.6	2.640	
GJR-GARCH(1,1)	6,983.7	7,014.6	2.669	
EGARCH(1,1)	—	—	—	Failed*

* Analytic forecasts not available for $h > 1$

Result

GARCH(1,1) selected based on lowest validation MAE for volatility forecasts.

Procedure

Refit GARCH(1,1) on Training + Validation, evaluate on **held-out test set** using **rolling one-step-ahead forecasts**.

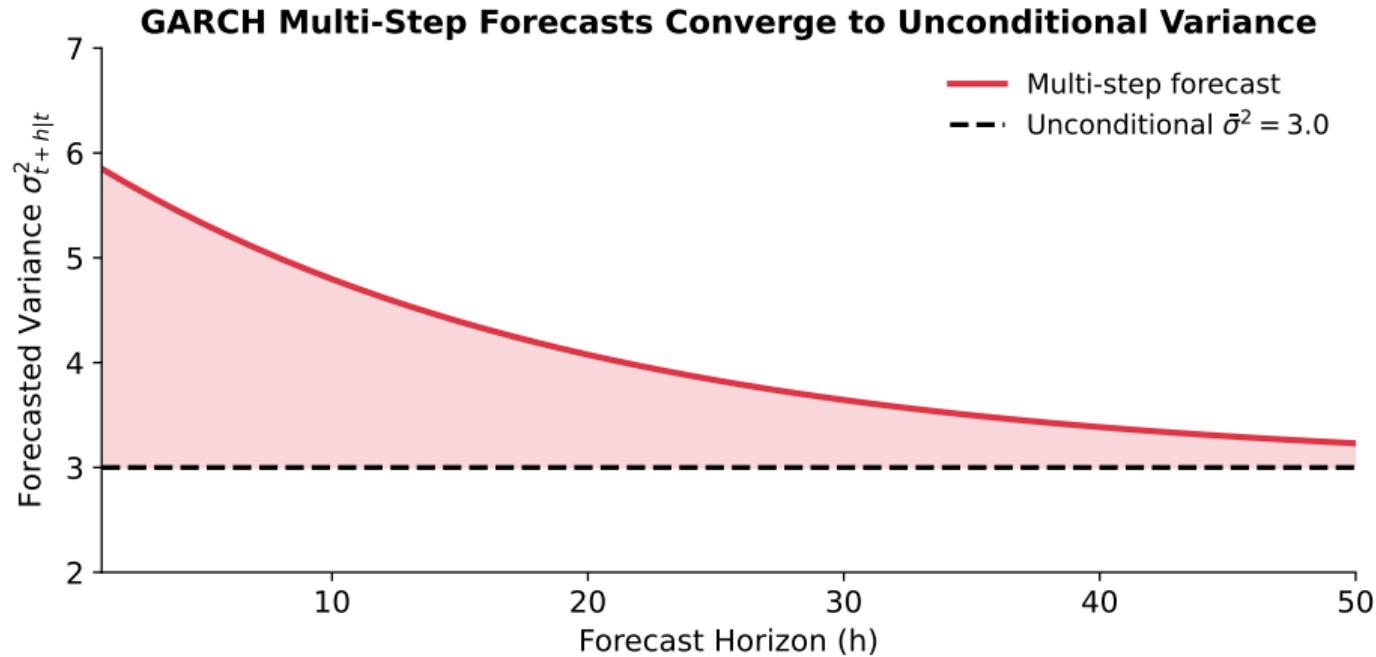
Estimated Parameters		
Param	Estimate	Std Err
ω	0.239	0.088
α_1	0.120	0.021
β_1	0.879	0.020
$\alpha_1 + \beta_1$	0.999	

Test Set Performance	Metric	Value
	Volatility MAE	1.88
	Volatility RMSE	2.21

Interpretation

High persistence ($\alpha + \beta \approx 1$) confirms volatility clustering.

GARCH: Multi-Step Forecasts Converge



The Problem

GARCH forecasts converge to unconditional variance

Prof. Daniel Traian Pele, PhD

Intuition

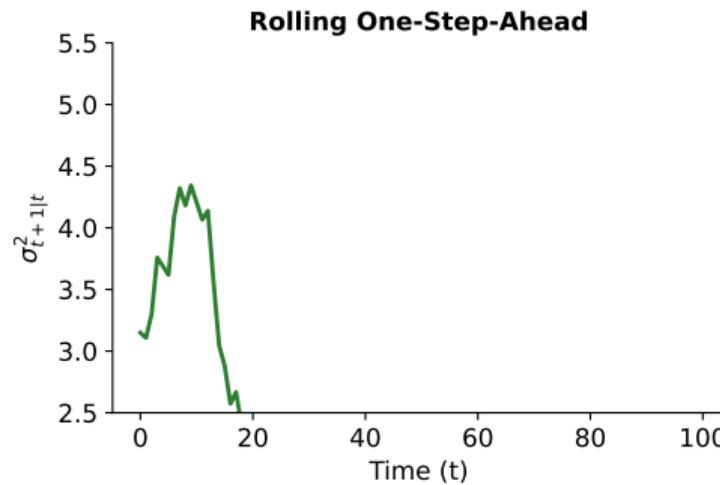
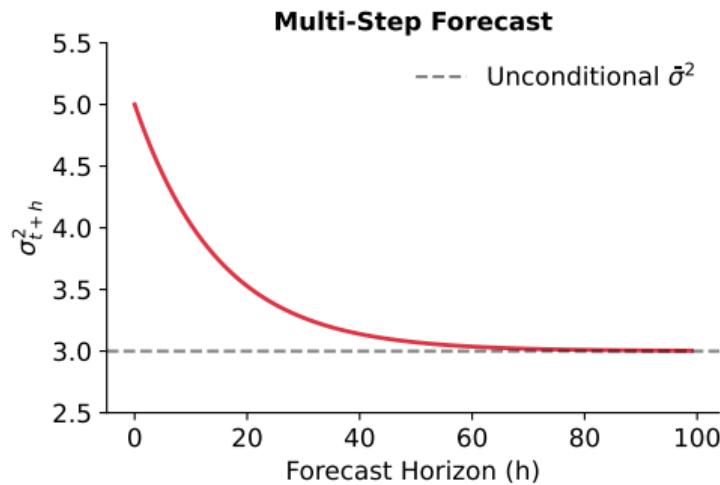
After a shock, volatility gradually returns to its long-run

Chapter 10: Comprehensive Review

Academic Year 2025–2026

13 / 37

GARCH: Rolling One-Step-Ahead Solution



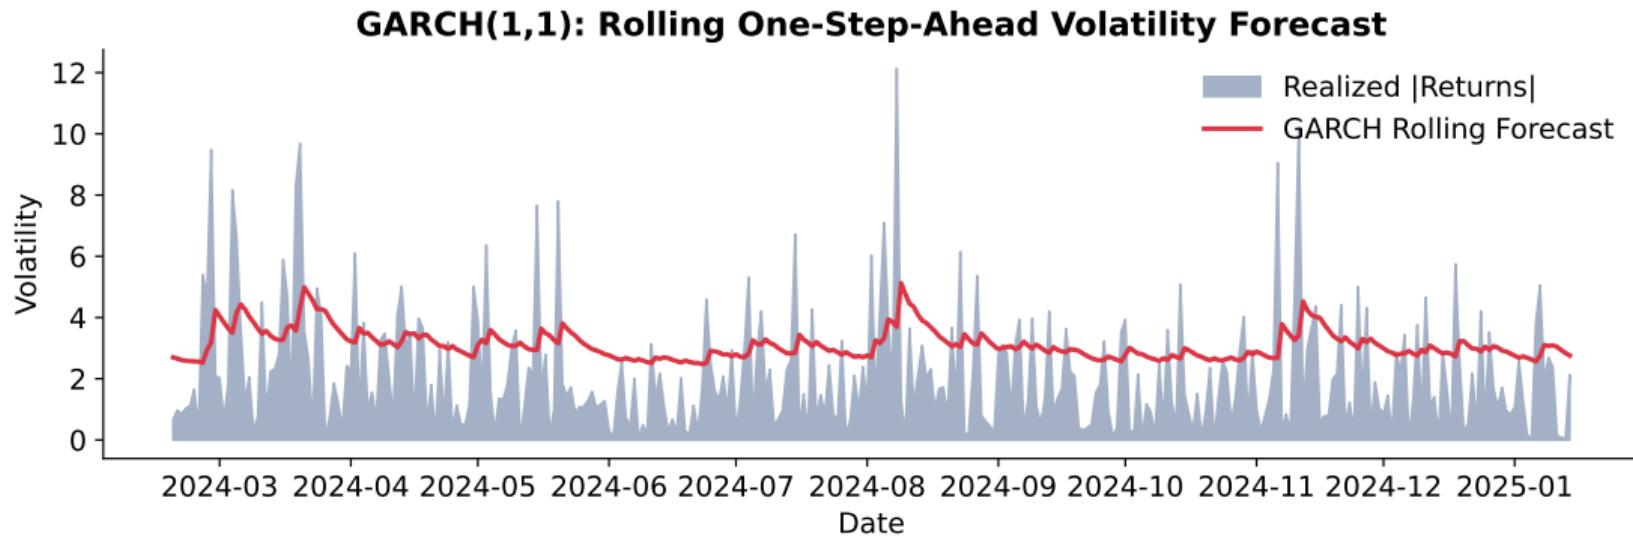
Multi-Step (Left)

Converges to $\bar{\sigma}^2$ (flat)

Rolling 1-Step (Right)

Re-estimate at each t (dynamic)

Bitcoin: GARCH Volatility Forecast (Test Set)



Result

Rolling one-step-ahead GARCH(1,1) forecasts capture the **dynamic volatility patterns**. The forecast (red line) tracks the realized volatility (blue area), demonstrating the predictability of variance.

Summary

- ❶ Returns are stationary; prices are not
- ❷ GARCH(1,1) outperforms more complex variants
- ❸ High persistence ($\alpha + \beta = 0.999$)
- ❹ Volatility is predictable even when returns are not

Limitations

- GARCH assumes **symmetric** shocks
- Does not capture **jumps**
- Normal distribution may be restrictive

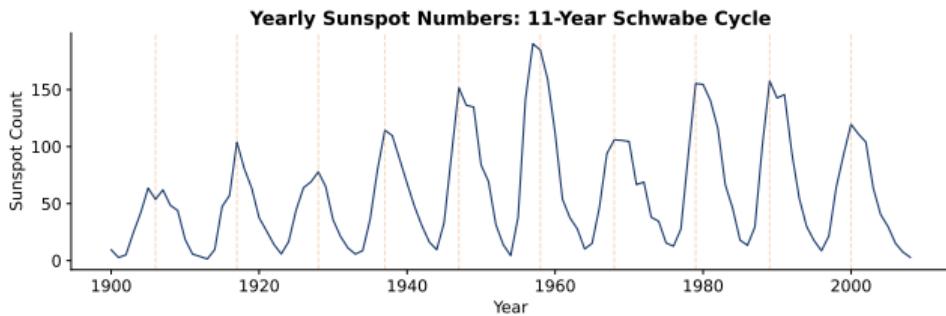
Practical Implications

- Risk management: VaR, Expected Shortfall
- Option pricing requires volatility forecasts
- Portfolio optimization with time-varying risk

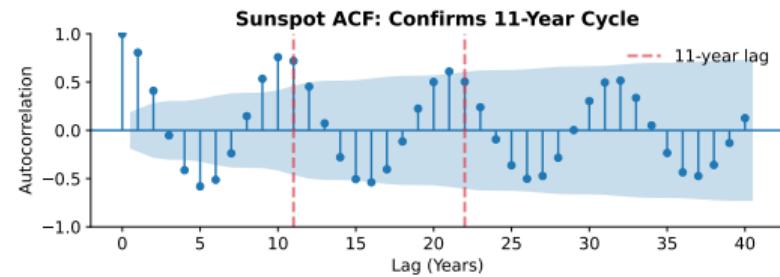
Extensions

- Student-t innovations
- Realized volatility
- HAR models

Sunspots: The 11-Year Solar Cycle



Dashed lines mark cycle peaks (\approx every 11 years). Amplitude varies.



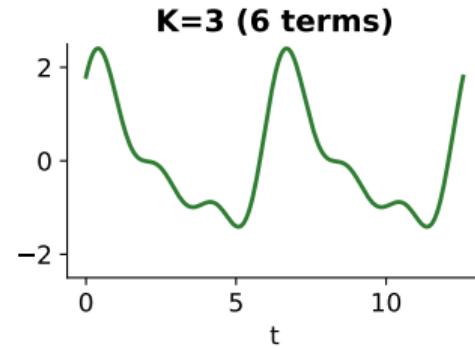
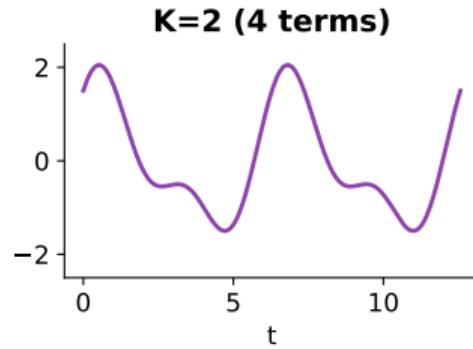
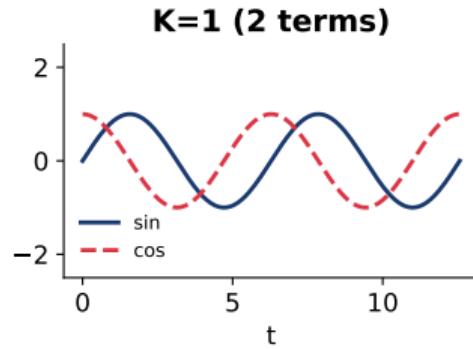
ACF peaks at lag 11 and 22 confirm the solar cycle periodicity.

Challenge

SARIMA(p, d, q)(P, D, Q)₁₁ requires estimating seasonal lags at 11, 22, 33... Too many parameters! **Solution:** Use Fourier terms instead.

Fourier Terms for Seasonality

Fourier Terms: More K = More Flexibility



How It Works

Approximate any periodic pattern using sine and cosine waves:

$$S_t = \sum_{k=1}^K \left[\alpha_k \sin\left(\frac{2\pi k t}{s}\right) + \beta_k \cos\left(\frac{2\pi k t}{s}\right) \right]$$

Key Insight

- $K = 1$: Simple wave (2 params)
- $K = 3$: Complex shape (6 params)
- For sunspots: $s = 11$, $K = 3$

Sunspots: Model Selection

Methodology

Compare $K = 1, 2, 3, 4$ Fourier harmonics on validation set.

	Set	Period	N
Data Split	Training	1900–1975	76
	Validation	1976–1991	16
	Test	1992–2008	17
Total			109

	K	AIC	Val RMSE
Model Comparison	1	665.9	87.15
	2	668.0	86.92
	3	671.8	86.81
	4	674.5	87.93

Result

$K = 3$ Fourier harmonics selected (6 parameters for 11-year cycle).

Sunspots: Test Set Results

Final Model

ARIMA(2,0,1) + 3 Fourier harmonics

Significant Coefficients:

Term	Coef	p-value
\sin_1	34.71	< 0.001
\cos_1	-29.21	0.018
AR(1)	1.34	< 0.001

Test Performance

Metric	Value
RMSE	48.51
MAE	39.31

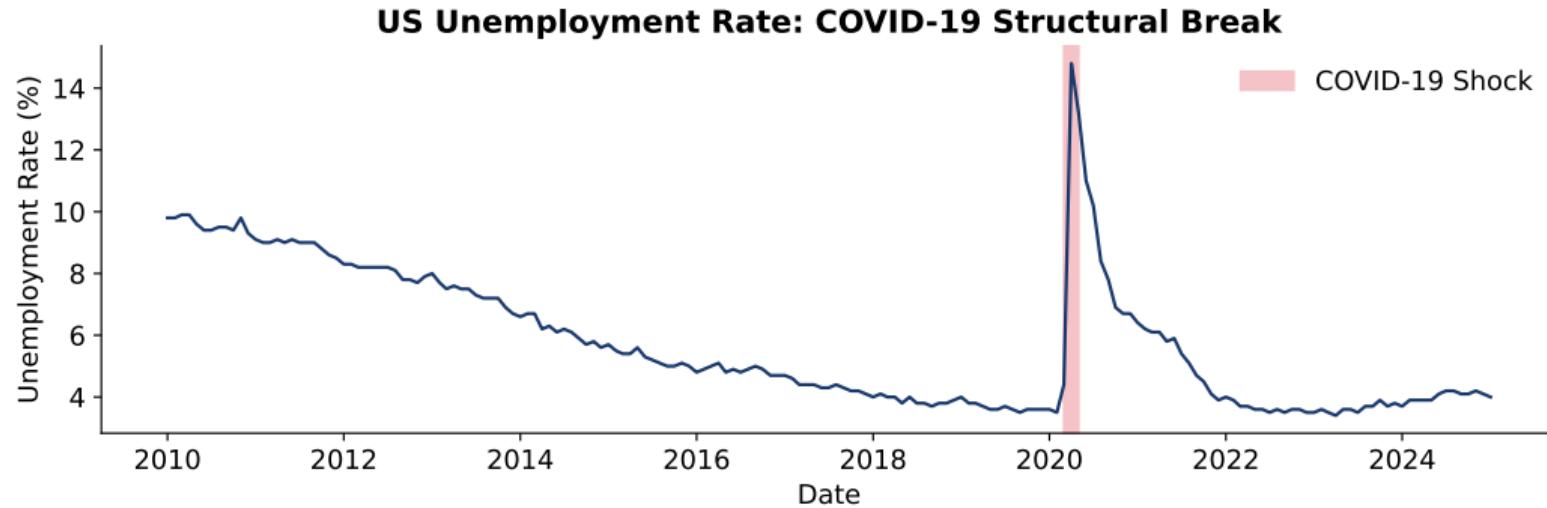
Note

High MAPE due to near-zero values at solar minimum.

Key Insight

Fourier terms efficiently capture the 11-year cycle with only 6 parameters.

Unemployment: COVID-19 Structural Break



The Challenge

- Pre-COVID: 3.5% (50-year low)
- April 2020: **14.8%** peak
- Largest monthly jump in US history

Why Not ARIMA?

ARIMA treats sudden jumps as **outliers**. Prophet detects **changepoints** and adapts the trend accordingly.

Definition 3 (Prophet Decomposition)

Prophet models time series as:

$$y_t = g(t) + s(t) + h(t) + \varepsilon_t \quad (4)$$

- $g(t)$: Piecewise linear/logistic **trend** with changepoints
- $s(t)$: Fourier-based **seasonality**
- $h(t)$: **Holiday** effects
- ε_t : Error term

Changepoint Detection

- Automatic selection of changepoint locations
- `changepoint_prior_scale` controls flexibility
- Higher = more changepoints

Advantages

- Handles missing data
- Interpretable components
- Robust to outliers
- Uncertainty quantification

Hyperparameter Tuning

Tune `changepoint_prior_scale` on validation set.

	Set	Period	N
Data Split	Training	2010-01 to 2019-09	117
	Validation	2019-10 to 2021-10	25
	Test	2021-11 to 2025-01	38
Total			180

	Scale	Val RMSE	
	0.01	4.21	
	0.05	3.89	
Scale Comparison	0.10	3.52	Best
	0.30	3.67	
	0.50	3.81	

Interpretation

Scale = 0.10 balances flexibility (capturing COVID shock) with stability.

Unemployment: Results

Test Set Performance

Metric	Value
RMSE	0.42
MAE	0.35
MAPE	9.2%

Detected Changepoints

- 2020-03: COVID onset
- 2020-05: Recovery begins
- 2022-01: Stabilization

Key Finding

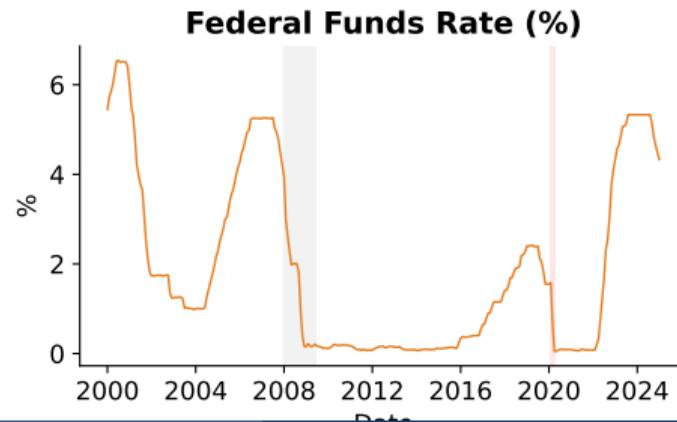
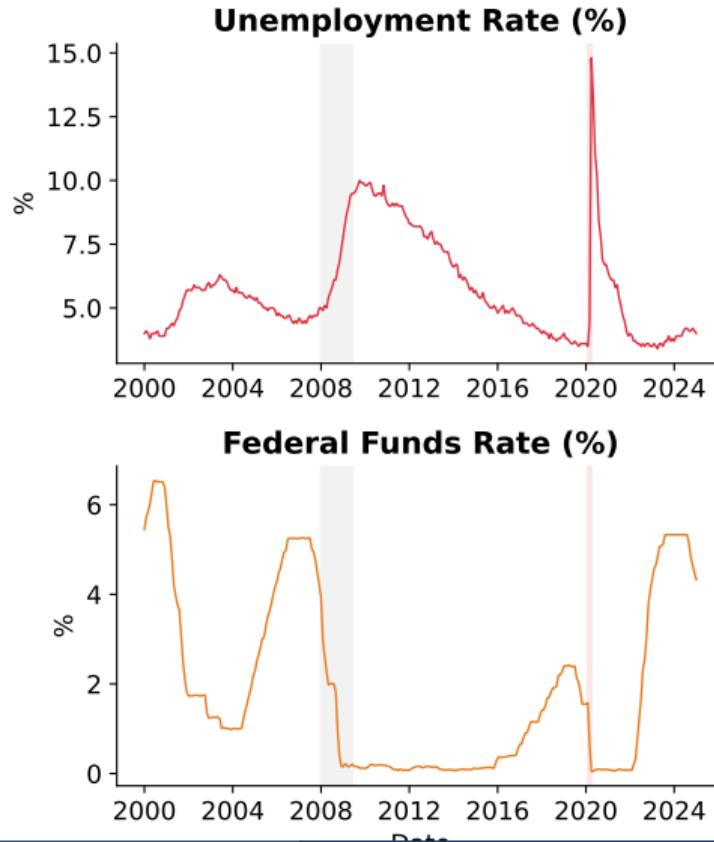
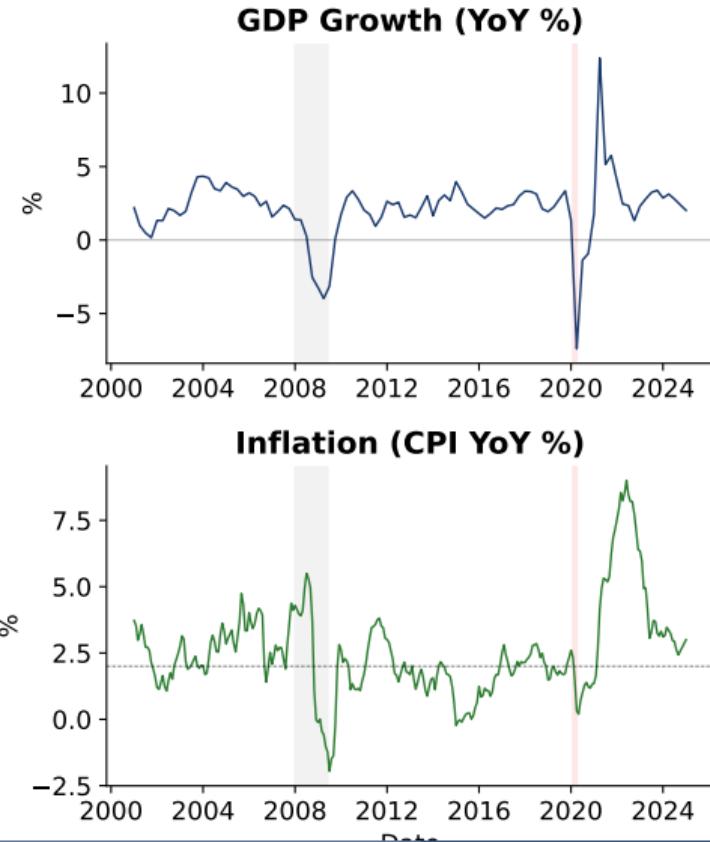
Prophet successfully:

- Detected COVID changepoint
- Adapted trend post-shock
- Provided uncertainty bands

Practical Value

- Economic policy analysis
- Labor market monitoring
- Early warning system

VAR: Multivariate Economic Data



Definition 4 (Vector Autoregression VAR(p))

For K variables $y_t = (y_{1t}, \dots, y_{Kt})'$:

$$y_t = c + A_1 y_{t-1} + A_2 y_{t-2} + \dots + A_p y_{t-p} + u_t \quad (5)$$

where A_i are $K \times K$ coefficient matrices and $u_t \sim N(0, \sigma^2)$.

For Our 4-Variable System

VAR(2) has:

- 4 intercepts
- $2 \times 4 \times 4 = 32$ AR coefficients
- **36 parameters total**

Lag Selection

Use information criteria:

- **AIC**: Tends to overfit
- **BIC**: More parsimonious
- Cross-validation on held-out data

VAR: Lag Selection and Estimation

Information Criteria

Lag	BIC
1	-4.810
2	-5.178
3	-4.633
4	-4.614

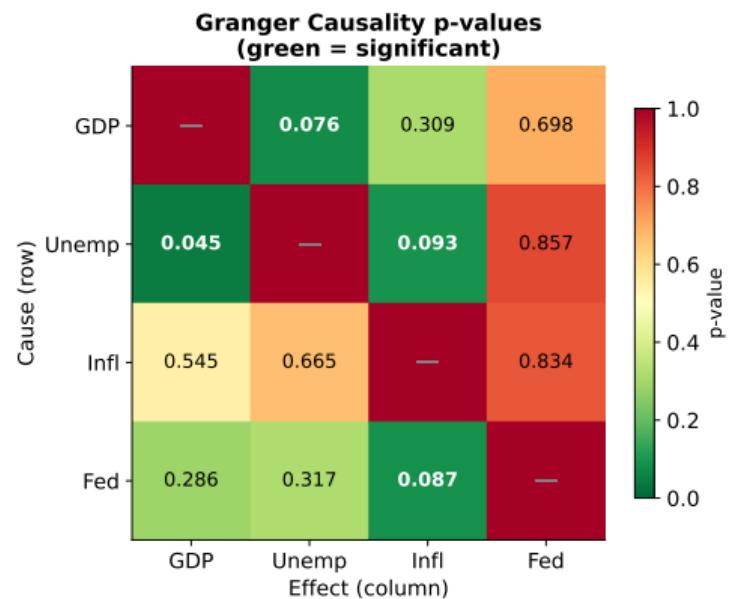
Data Split

Set	Period	N
Training	2001-Q1 to 2017-Q4	68
Validation	2018-Q1 to 2021-Q2	14
Test	2021-Q3 to 2024-Q3	14
Total		96

Validation Check

VAR(2) also achieves lowest validation RMSE.

Granger Causality Analysis



What is Granger Causality?

X Granger-causes Y if past X improves prediction of Y beyond past Y alone.

Warning: “Granger causality” \neq true causality!

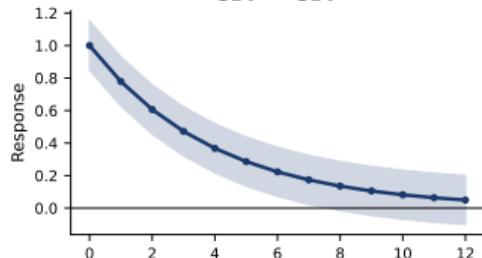
Economic Findings

- Unemp \rightarrow GDP ($p = 0.045$): Okun's Law
- Fed \rightarrow Inflation ($p = 0.087$): Monetary policy works

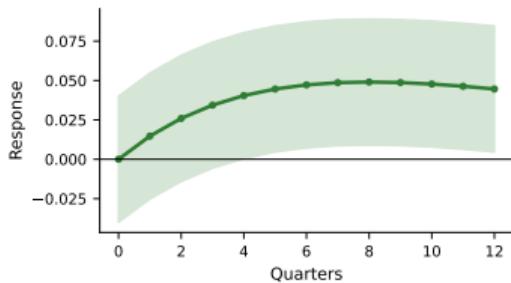
Green cells: $p < 0.10$ (significant). Read: row causes column.

Impulse Response Functions (IRF)

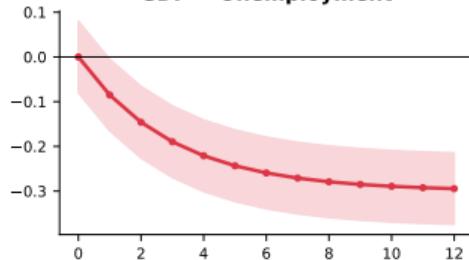
Impulse Response Functions: Response to GDP Shock
 $GDP \rightarrow GDP$



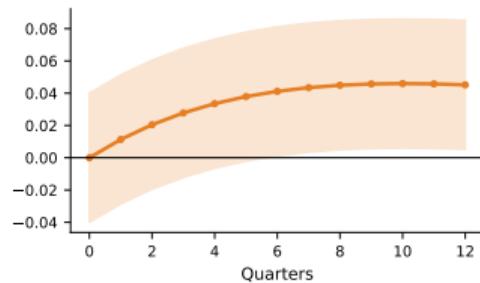
$GDP \rightarrow Inflation$



Impulse Response Functions: Response to GDP Shock
 $GDP \rightarrow Unemployment$



$GDP \rightarrow Fed Rate$



What is IRF?

Shows how a 1-unit shock to one variable affects others over time.

GDP Shock Effects

- **Unemp** ↓: Okun's Law
- **Inflation** ↑: Demand-pull
- **Fed Rate** ↑: Taylor Rule

Test Set Performance by Variable

Variable	RMSE	MAE	Direction Acc.
GDP Growth	2.18	1.72	71%
Unemployment	0.89	0.71	79%
Inflation	1.24	0.98	64%
Fed Rate	0.95	0.78	71%
Average	1.32	1.05	71%

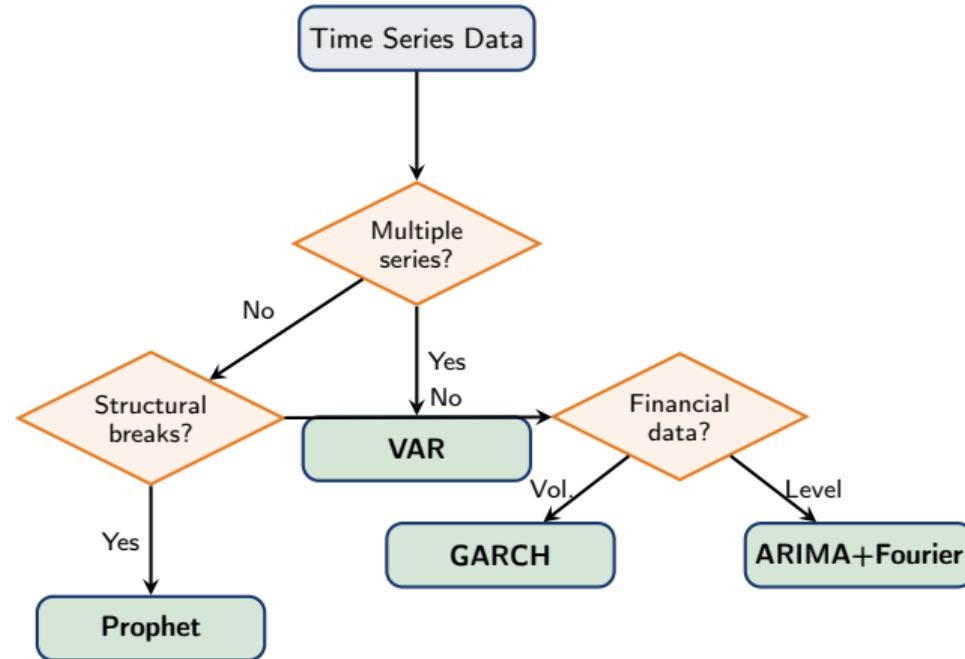
Strengths

- Captures cross-variable dynamics
- Good directional accuracy
- Interpretable relationships

Limitations

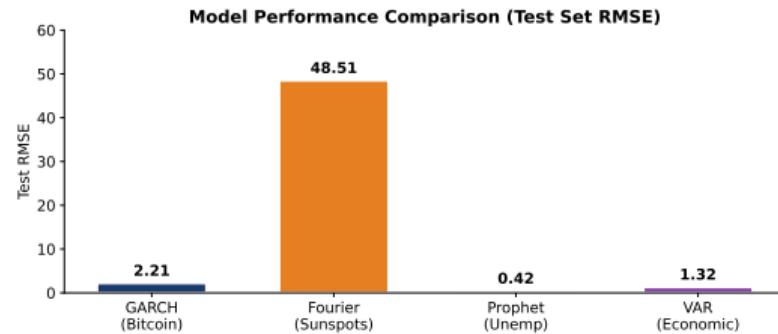
- Many parameters (curse of dimensionality)
- Sensitive to lag selection
- COVID period challenging

Model Selection Framework



Summary: Model Comparison

Case	Challenge	Model	RMSE
Bitcoin	Volatility	GARCH	2.21
Sunspots	Seasonality	Fourier	48.51
Unemp	Break	Prophet	0.42
Economic	Multi-var	VAR	1.32



Key Principle

Match the model to the data characteristics. No single model dominates—choose based on:

- Nature of the forecasting problem (level vs. volatility)
- Data properties (seasonality, breaks, multiple series)
- Interpretability requirements

Best Practices for Applied Forecasting

Methodology

- ① Explore data thoroughly
- ② Test for stationarity
- ③ Split train/validation/test
- ④ Compare models on validation
- ⑤ Report test set metrics

Practical Tips

- Start simple (random walk, naive)
- Add complexity only if needed
- Visualize forecasts vs actuals
- Check residuals for patterns
- Report confidence intervals

Common Mistakes

- Peeking at test data
- Over-fitting to training set
- Ignoring model assumptions
- Not reporting uncertainty

Remember

"All models are wrong, but some are useful."
— George E. P. Box

Key Takeaways

① Rigorous Methodology

- Train/validation/test split prevents overfitting
- Test set must remain untouched until final evaluation

② Match Model to Data

- Financial volatility → GARCH
- Long seasonality → Fourier terms
- Structural breaks → Prophet
- Multiple series → VAR

③ Interpret Results Carefully

- Granger causality \neq true causality
- Out-of-sample performance matters most
- Simpler models often work better

References

-  Box, G.E.P., Jenkins, G.M., Reinsel, G.C., & Ljung, G.M. (2015). *Time Series Analysis: Forecasting and Control*. 5th ed., Wiley.
-  Hamilton, J.D. (1994). *Time Series Analysis*. Princeton University Press.
-  Tsay, R.S. (2010). *Analysis of Financial Time Series*. 3rd ed., Wiley.
-  Hyndman, R.J., & Athanasopoulos, G. (2021). *Forecasting: Principles and Practice*. 3rd ed., OTexts.
-  Taylor, S.J., & Letham, B. (2018). Forecasting at Scale. *The American Statistician*, 72(1), 37-45.
-  Bollerslev, T. (1986). Generalized Autoregressive Conditional Heteroskedasticity. *Journal of Econometrics*, 31(3), 307-327.
-  Sims, C.A. (1980). Macroeconomics and Reality. *Econometrica*, 48(1), 1-48.

Real Data Used in This Chapter

- **Bitcoin:** Yahoo Finance (BTC-USD), 2019–2025
- **Sunspots:** Statsmodels Wolfer dataset, 1900–2008
- **US Unemployment:** Federal Reserve FRED (UNRATE), 2010–2025
- **Economic Variables:** FRED (GDPC1, UNRATE, CPIAUCSL, FEDFUNDS), 2000–2025

Reproducibility

All analyses can be reproduced using the accompanying Jupyter notebook:
`chapter10_lecture_notebook.ipynb`

Thank You

Questions?

Prof. Daniel Traian Pele, PhD

danpele@ase.ro

Bucharest University of Economic Studies