



Time Series Analysis

Seminar 8: Long memory and ARFIMA



Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Acronyms used in this material

- ▣ ACF: Autocorrelation Function
- ▣ AI: Artificial Intelligence
- ▣ AIC: Akaike Information Criterion
- ▣ AR: AutoRegressive (model); in event studies: Abnormal Return
- ▣ ARCH: AutoRegressive Conditional Heteroskedasticity
- ▣ ARFIMA: AutoRegressive Fractionally Integrated Moving Average (model)
- ▣ ARIMA: AutoRegressive Integrated Moving Average (model)
- ▣ ARMA: AutoRegressive Moving Average (model)
- ▣ BET: Bucharest Exchange Trading (index)
- ▣ BIC: Bayesian Information Criterion
- ▣ BNR: Banca Națională a României (the National Bank of Romania)
- ▣ CPI: Consumer Price Index
- ▣ CPIAUCSL: Consumer Price Index for All Urban Consumers (FRED series code)
- ▣ DFA: Detrended Fluctuation Analysis
- ▣ EODHD: EOD Historical Data (market-data provider)
- ▣ FIGARCH: Fractionally Integrated GARCH
- ▣ FRED: Federal Reserve Economic Data
- ▣ GARCH: Generalised AutoRegressive Conditional Heteroskedasticity
- ▣ GPH: Geweke–Porter–Hudak (estimator)
- ▣ HICP: Harmonised Index of Consumer Prices
- ▣ LW: Local Whittle (estimator)
- ▣ ML: Machine Learning
- ▣ OLS: Ordinary Least Squares
- ▣ RV: Realised Volatility
- ▣ SE: Standard Error
- ▣ US: United States

Today's question and route

- **Question:** does a shock to a series fade at an exponential rate (ARMA) or much more slowly (long memory), and how do we measure the difference?
 - ▶ this seminar comes **before** Lecture 8: the section “What you need today” gives every definition the tasks use
- Route
 - ▶ Part A: fractional weights and ARFIMA autocorrelations; R/S and GPH by hand; an ARFIMA forecast; GARCH and FIGARCH weights
 - ▶ Part B: the Nile and its break; Romanian inflation; S&P 500 volatility; BET and EUR/RON
 - ▶ Part C: US inflation and monetary regimes; an AI answer to audit
- Notebook for today: [open the seminar notebook in Google Colab](#); each task names its notebook section

Exercise map

Task	Question	Type	Model
A1, A2	weights of $(1 - L)^d$, ACF of ARFIMA(0, d , 0), ranges of d	Solved, Proposed	A1
A3, A4	R/S and H by hand; a GPH slope by hand	Solved, Proposed	A3
A5, A6	an ARFIMA forecast; GARCH and FIGARCH weights	Solved, Proposed	A5
B1, B2	the Nile and its break; Romanian inflation	Solved, Proposed	B1
B3, B4	S&P 500 volatility; BET and EUR/RON	Solved, Proposed	B3
C1, C2	US inflation and regimes; what is wrong in an AI answer?	Proposed	B1, B2

- **[Solved]**: full solution in the slides and in the notebook, a model to follow; **[Proposed]**: you solve it, following the model

Data used

Series	Source	Frequency	Period
Nile flow at Aswan	statsmodels (nile)	yearly	1871–1970
HICP, Romania	Eurostat (prc_hicp_minr)	monthly, 2015 = 100	2005–2026
CPI, United States	FRED (CPIAUCSL)	monthly, SA	1947–2026
S&P 500, BET	EODHD	daily, close	2000–2026
EUR/RON	BNR reference rate	daily	2005–2026

- Romanian inflation: monthly rate $100\Delta \ln P_t$ minus its monthly means (seasonal adjustment), and the 12-month rate $100(\ln P_t - \ln P_{t-12})$
- In the notebook: `gph`, `local_whittle`, `hurst_rs`, `arfima_exact_ml`, `monthly_rv`; no account or key is needed

What you need today (1/4): short and long memory

- ▣ **Short memory** (ARMA, [Chapter 2](#)): $|\rho(k)| \leq Cr^k$, $0 < r < 1$; the autocorrelations are summable
 - ▶ AR(1): $\rho(k) = \phi^k$
- ▣ **Long memory**: $\rho(k) \sim Ck^{2d-1}$, $0 < d < 1/2$: hyperbolic decay, $\sum_k \rho(k) = \infty$
 - ▶ equivalently, the spectral density has a pole at zero: $f(\lambda) \sim G\lambda^{-2d}$ as $\lambda \rightarrow 0$
- ▣ **Hurst exponent** $H = d + 1/2$: $H = 0.5$ no memory, $H > 0.5$ persistence, $H < 0.5$ anti-persistence
 - ▶ Hurst found $H \approx 0.7$ for the Nile [Hurst \(1951\)](#)

What you need today (2/4): fractional differencing and ARFIMA

- $(1 - L)^d = \sum_k \pi_k L^k$ with $\pi_0 = 1$, $\pi_k = \pi_{k-1}(k - 1 - d)/k$; $(1 - L)^{-d}$: $\psi_k = \psi_{k-1}(k - 1 + d)/k$
 - ▶ **ARFIMA**(p, d, q): $\phi(L)(1 - L)^d(x_t - \mu) = \theta(L)\varepsilon_t$ Granger and Joyeux (1980), Hosking (1981)
- Ranges: invertible for $d > -1/2$; stationary for $d < 1/2$; mean-reverting for $d < 1$; $d = 1$ unit root
 - ▶ if $1/2 \leq d < 3/2$, difference once: $(1 - L)x_t$ is ARFIMA with $d - 1$
- ARFIMA(0, d , 0): $\rho(1) = d/(1 - d)$, $\rho(k) = \rho(k - 1)(k - 1 + d)/(k - d)$, $\gamma(0) = \sigma^2\Gamma(1 - 2d)/\Gamma(1 - d)^2$
 - ▶ one-step forecast: $\hat{x}_{T+1} = \mu - \sum_{k \geq 1} \pi_k(x_{T+1-k} - \mu)$

What you need today (3/4): estimators of d and H

- ▣ **R/S**: in a block, $Y_k = \sum_{i \leq k} (x_i - \bar{x})$, $R = \max Y_k - \min Y_k$, S the standard deviation (divisor n)
 - ▶ $H = \text{slope of } \log_{10}(R/S)_n \text{ on } \log_{10} n$; **DFA**: slope of the detrended fluctuation $\log F(n)$
- ▣ **GPH**: OLS of $\log I(\lambda_j)$ on $X_j = -\log(4 \sin^2(\lambda_j/2))$, $j = 1..m$; slope $= \hat{d}$, $SE = \pi/\sqrt{24m}$
 - ▶ $I(\lambda_j)$: periodogram at $\lambda_j = 2\pi j/T$; OLS slope $= \sum (X_j - \bar{X})(y_j - \bar{y}) / \sum (X_j - \bar{X})^2$
- ▣ **Local Whittle Robinson (1995)**: a local likelihood on the same m frequencies, $SE = 1/(2\sqrt{m})$; bandwidth $m = \lfloor T^{0.65} \rfloor$
 - ▶ **Exact ML Sowell (1992)**: the full ARFIMA(p, d, q) likelihood; compare models by AIC and BIC

What you need today (4/4): volatility and spurious memory

- Returns r_t are almost uncorrelated; $|r_t|$ and realised volatility have long memory [Ding, Granger and Engle \(1993\)](#)
 - ▶ monthly **realised volatility**: $\sqrt{RV_m}$, $RV_m = \sum_{t \in m} r_t^2$; we model its logarithm
 - ▶ **shuffle test**: permute the days at random; real memory disappears, the distribution stays
- GARCH(1,1) as ARCH(∞): weight $\alpha\beta^{k-1}$ on ε_{t-k}^2 ; FIGARCH [Baillie, Bollerslev and Mikkelsen \(1996\)](#): weights $\approx k^{-1-d}$
 - ▶ with $\phi = \beta = 0$, FIGARCH weights are $\lambda_k = -\pi_k$ of $(1-L)^d$
- **Spurious long memory**: level shifts and rare regime switches give $\hat{d} > 0$ for short-memory series [Diebold and Inoue \(2001\)](#); re-estimate after removing the breaks

A1: fractional weights and the ACF [Solved]

Task

- ▣ A series is ARFIMA(0, 0.25, 0) with $\sigma^2 = 1$.
- ▣ 1. Compute $\pi_1, \dots, \pi_4$ of $(1 - L)^{0.25}$.
- ▣ 2. Compute $\rho(1)$ and $\rho(2)$, and give H and the class of the process.
- ▣ 3. Compare $\rho(5)$ with that of an AR(1) with the same $\rho(1)$.
- ▣ Report: four weights, two autocorrelations, H , two values of $\rho(5)$.

Solution

- ▣ 1. $\pi_1 = -0.25$; $\pi_2 = -0.25 \cdot 0.75/2 = -0.0938$;
 $\pi_3 = -0.0938 \cdot 1.75/3 = -0.0547$;
 $\pi_4 = -0.0547 \cdot 2.75/4 = -0.0376$
- ▣ 2. $\rho(1) = 0.25/0.75 = 0.333$; $\rho(2) = 0.333 \cdot 1.25/1.75 = 0.238$;
 $H = 0.75$: stationary, long memory
- ▣ 3. ARFIMA: $\rho(5) = 0.151$; AR(1): $0.333^5 = 0.0041$, about 37 times smaller

A2: which range is d in? [Proposed]

Solution discussed in class

Task

- ▣ Two series: (a) ARFIMA(0, 0.6, 0); (b) ARFIMA(0, -0.3, 0). Model: A1.
- ▣ 1. Compute π_1 and π_2 of $(1 - L)^d$ for each.
- ▣ 2. Say whether each series is stationary, invertible and mean-reverting.
- ▣ 3. For (b) compute $\rho(1)$ and $\rho(2)$; for (a) compute $\rho(1)$ of the first difference.
- ▣ Report: four weights, two classifications, three autocorrelations.

A3: R/S by hand and H from three points [Solved]

Task

- ▣ Block: $x = (0.6, 0.4, 0.5, 0.3, -0.4, -0.6, -0.2, -0.5)$. An R/S plot gives $(R/S)_n = 4.1; 9.6; 22.4$ for $n = 16, 64, 256$.
- ▣ 1. Compute $\bar{x}$, the cumulative deviations Y_k , R , S and R/S .
- ▣ 2. Compute H as the slope of $\log_{10}(R/S)_n$ on $\log_{10} n$, and $d = H - 0.5$.
- ▣ Report: R/S , H , d and one sentence.

Solution

- ▣ 1. $\bar{x} = 0.0125$; Y_k : 0.5875, 0.9750, 1.4625, 1.7500, 1.3375, 0.7250, 0.5125, 0; $R = 1.750$, $S = 0.457$, $R/S = 3.83$
- ▣ the signs are grouped (four positive, then four negative): the range is large, as for a persistent series
- ▣ 2. $\log_{10} n$: 1.204, 1.806, 2.408; $\log_{10}(R/S)$: 0.613, 0.982, 1.350; slope $H = 0.61$, $d = 0.11$: moderate persistence

A4: a GPH slope by hand [Proposed]

Solution discussed in class

Task

- ▣ For $m = 4$ frequencies: $X_j = 0.5; 1.6; 2.3; 3.0$ and $\log I(\lambda_j) = 1.2; 1.8; 2.0; 2.5$. Model: A3.
- ▣ 1. Compute $\bar{X}$, $\bar{y}$, $\sum(X_j - \bar{X})(y_j - \bar{y})$ and $\sum(X_j - \bar{X})^2$.
- ▣ 2. Compute $\hat{d}$ and its standard error $\pi/\sqrt{24m}$.
- ▣ 3. Test $H_0: d = 0$ at 5% and comment on the size of m .
- ▣ Report: four sums, $\hat{d}$, its SE, a decision and one sentence.



A5: a one-step ARFIMA forecast [Solved]

Task

- ▣ ARFIMA(0, 0.4, 0) with $\mu = 2$; the last four values: $x_T = 3.0$, $x_{T-1} = 2.5$, $x_{T-2} = 2.8$, $x_{T-3} = 1.5$.
- ▣ 1. Write the first four weights $-\pi_k$ of $(1 - L)^{0.4}$.
- ▣ 2. Compute $\hat{x}_{T+1}$ using these four lags.
- ▣ 3. Compare with an AR(1) with the same $\rho(1)$.
- ▣ Report: four weights, two forecasts and one sentence.

Solution

- ▣ 1. $-\pi_k$: 0.4; 0.12; 0.064; 0.0416
- ▣ 2.

$$\hat{x}_{T+1} = 2 + 0.4(1.0) + 0.12(0.5) + 0.064(0.8) + 0.0416(-0.5) = 2 + 0.4000 + 0.0600 + 0.0512 - 0.0208 = 2.490$$
- ▣ 3. AR(1): $\phi = \rho(1) = 0.4/0.6 = 0.667$,
 $\hat{x}_{T+1} = 2 + 0.667 \cdot 1.0 = 2.667$; the AR(1) looks only at x_T ,
 ARFIMA spreads the weight over the past (in practice over all T values)

A6: how long does a volatility shock count? [Proposed]

Solution discussed in class

Task


- ▣ GARCH(1,1) with $\alpha = 0.10$, $\beta = 0.88$; FIGARCH(0, d , 0) with $d = 0.4$ (weights $\lambda_k = -\pi_k$). Model: A5.
- ▣ 1. Compute the weight of ε_{t-k}^2 in σ_t^2 for $k = 1, 2, 3$ in both models.
- ▣ 2. Compute both weights at $k = 20$ and $k = 100$; for FIGARCH use $\lambda_k \approx d k^{-1-d} / \Gamma(1-d)$, $\Gamma(0.6) = 1.489$.
- ▣ 3. Explain which model keeps a crisis in today's volatility longer.
- ▣ Report: ten weights and one sentence.



B1: does the Nile have long memory? [Solved]

- ▣ **Question:** is the persistence of the Nile flow long memory, or the effect of the change in 1898?
- ▣ **Inputs:** yearly flow at Aswan, 1871–1970 ($T = 100$)
- ▣ **Tasks**
 1. Plot the series and its ACF up to lag 20.
 2. Estimate H by R/S, d by GPH and local Whittle ($m = \lfloor T^{0.65} \rfloor$) and by exact ML of ARFIMA(0, d , 0).
 3. Subtract the mean of 1871–1898 and the mean of 1899–1970 and repeat step 2.
 4. Interpretation: does the Nile have long memory?
- ▣ **Report:** two ACF values, four estimates before and after, one sentence
- ▣ **Notebook:** section B1

B1: solution [Solved]



- Raw: $\hat{\rho}(1) = 0.50$, $\hat{\rho}(10) = 0.09$; $H_{R/S} = 0.78$; GPH 0.44 (SE 0.15), LW 0.40 (SE 0.11), ML 0.36 (SE 0.07); $m = 19$
- Means 1098 and 850; after removing them: $\hat{\rho}(1) = 0.16$; $H_{R/S} = 0.63$; GPH -0.07, LW -0.17, ML 0.05 (SE 0.10)
- Interpretation: for 1871–1970 the evidence of long memory comes from one break; R/S stays above 0.5 because it is biased upwards in short samples


B2: long memory in Romanian inflation [Proposed]

- **Question:** how persistent is Romanian inflation, and does the way we measure inflation change the answer?
- **Inputs:** monthly HICP, 2005–2026: the monthly rate seasonally adjusted with the monthly means, and the 12-month rate; model: B1
- **Tasks**
 1. Estimate d of the monthly rate by GPH and local Whittle with $m = \lfloor T^a \rfloor$, $a = 0.5; 0.65; 0.8$.
 2. Fit ARFIMA(0, d , 0) and ARMA(1,1) by exact ML and compare their BIC.
 3. Estimate d of the 12-month rate by local Whittle and compare the two ACFs.
 4. Interpretation: why does the 12-month rate give d close to 1?
- **Report:** six estimates, two BIC values, one d , two sentences
- **Notebook:** section B2

B3: memory in S&P 500 volatility [Solved]

- **Question:** where is the memory of stock returns: in the returns or in their size?
- **Inputs:** daily log returns of the S&P 500 since 2000 and their monthly realised volatility
- **Tasks**
 1. Estimate d of r_t and of $|r_t|$ by local Whittle and GPH.
 2. Shuffle the days at random and re-estimate d of $|r_t|$; plot both ACFs of $|r_t|$.
 3. Estimate d of the monthly log realised volatility.
 4. Interpretation: what does the shuffle test show?
- **Report:** eight estimates, the chart and one sentence
- **Notebook:** section B3

B3: solution [Solved]



- $T = 6714$ days; r_t : LW -0.01, GPH 0.02 (SE 0.03, 0.04); $|r_t|$: LW 0.53, GPH 0.56
- Shuffled $|r_t|$: LW -0.03, GPH -0.03; ACF of $|r_t|$: 0.289 (lag 1), 0.127 (lag 50), 0.009 (lag 200); log RV (321 months): LW 0.41, GPH 0.37
- Interpretation: the shuffle keeps every daily value and destroys their order; the memory disappears, so it lives in the clustering of calm and turbulent periods, not in the distribution of returns

B4: BET and EUR/RON [Proposed]

- **Question:** do the BET and the EUR/RON rate show the same memory, and is it stable before and after the 2008 crisis?
- **Inputs:** daily log returns of the BET (since 2000) and of the EUR/RON reference rate (since July 2005); model: B3
- **Tasks**
 1. Estimate d of r_t and $|r_t|$ by local Whittle on the full sample.
 2. Repeat on the subsamples up to 2007 and 2010–2026.
 3. Compare the estimates for r_t with the 95% Monte Carlo band of i.i.d. series of 1900 days.
 4. Interpretation: is the full-sample memory of BET returns real?
- **Report:** twelve estimates, one band and two sentences
- **Notebook:** section B4

C1: US inflation and monetary regimes [Proposed]

- ▣ **Question:** is the long memory of US inflation a property of inflation, or of the changes in monetary policy?
- ▣ **Inputs:** US monthly CPI inflation, annualised, 1947–2026; regimes 1947–1984, 1985–2019, 2020–2026; models: B1, B2
- ▣ **Tasks**
 1. Estimate d by local Whittle on the full sample and in each regime.
 2. Remove the three regime means and re-estimate d .
 3. Propose one more check (a break test, a Markov-switching model, a forecast comparison).
 4. Interpretation: can a constant d describe 80 years of inflation?
- ▣ **Report:** five estimates, one check and a project plan
- ▣ **Notebook:** section C1

C2: audit an AI answer [Proposed]

- ▣ A student asked an AI assistant about the long memory of Romanian inflation and of the S&P 500. The answer:
- ▣ (a) Monthly inflation has $d = 0.31$, so its Hurst exponent is $H = d - 0.5 < 0$: it is anti-persistent.
- ▣ (b) The 12-month rate gives $d = 1.00$, so Romanian inflation has a unit root.
- ▣ (c) With $d = 0.31$ the series is stationary and its shocks die out.
- ▣ (d) Use `statsmodels.tsa.arima.model.ARIMA(x, order=(0, 0.31, 0))` to fit the ARFIMA model.
- ▣ (e) S&P 500 returns have $d = -0.01$ and $|r|$ has $d = 0.53$, so returns are predictable.
- ▣ (f) A significant GPH estimate proves long memory.
- ▣ Tasks
 - ▶ 1. For each statement, say whether it is correct; if not, give the correct statement and, where possible, the correct number from the notebook (section C2).
 - ▶ 2. Report: a list of six verdicts with one line of justification each.

Key takeaways

- ▣ $(1 - L)^d$: weights by the recursion $\pi_k = \pi_{k-1}(k - 1 - d)/k$; ARFIMA(0, d , 0): $\rho(1) = d/(1 - d)$
- ▣ $H = d + 1/2$; R/S and GPH are slopes of log-log regressions
- ▣ Romanian monthly inflation: $d \approx 0.3$; S&P 500: memory in $|r_t|$, not in r_t
- ▣ Breaks can create long memory: the Nile, BET returns, US inflation
- ▣ An AI answer is a draft: check H against d , the library, and the data transformation

After the seminar

- ▣ Lecture 8 develops each topic of today: hyperbolic decay, fractional differencing, ARFIMA, the estimators, forecasting, volatility, spurious memory
- ▣ Try the [Proposed] tasks in the notebook
- ▣ C1 can grow into a team project: inflation memory and monetary regimes in Romania
- ▣ Reading: [Baillie \(1996\)](#); [Graves et al. \(2017\)](#); [Ding, Granger and Engle \(1993\)](#)
- ▣ **The seminar is for practice and is not graded; the solutions of [Proposed] tasks are discussed in class**

References

- ▣ Baillie, R. T. (1996). Long memory processes and fractional integration in econometrics. *Journal of Econometrics*, 73(1), 5–59. doi:10.1016/0304-4076(95)01732-1
- ▣ Baillie, R. T., Bollerslev, T., & Mikkelsen, H. O. (1996). Fractionally integrated generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 74(1), 3–30. doi:10.1016/S0304-4076(95)01749-6
- ▣ Diebold, F. X., & Inoue, A. (2001). Long memory and regime switching. *Journal of Econometrics*, 105(1), 131–159. doi:10.1016/S0304-4076(01)00073-2
- ▣ Ding, Z., Granger, C. W. J., & Engle, R. F. (1993). A long memory property of stock market returns and a new model. *Journal of Empirical Finance*, 1(1), 83–106. doi:10.1016/0927-5398(93)90006-D
- ▣ Granger, C. W. J., & Joyeux, R. (1980). An introduction to long-memory time series models and fractional differencing. *Journal of Time Series Analysis*, 1(1), 15–29. doi:10.1111/j.1467-9892.1980.tb00297.x
- ▣ Graves, T., Gramacy, R., Watkins, N., & Franzke, C. (2017). A brief history of long memory: Hurst, Mandelbrot and the road to ARFIMA, 1951–1980. *Entropy*, 19(9), 437. doi:10.3390/e19090437
- ▣ Hosking, J. R. M. (1981). Fractional differencing. *Biometrika*, 68(1), 165–176. doi:10.1093/biomet/68.1.165
- ▣ Hurst, H. E. (1951). Long-term storage capacity of reservoirs. *Transactions of the American Society of Civil Engineers*, 116(1), 770–799. doi:10.1061/TACEAT.0006518
- ▣ Robinson, P. M. (1995). Gaussian semiparametric estimation of long range dependence. *The Annals of Statistics*, 23(5), 1630–1661. doi:10.1214/aos/1176324317
- ▣ Sowell, F. (1992). Maximum likelihood estimation of stationary univariate fractionally integrated time series models. *Journal of Econometrics*, 53(1–3), 165–188. doi:10.1016/0304-4076(92)90084-5